BMI 702: Biomedical Artificial Intelligence

Foundations of Biomedical Informatics II, Spring 2023

Lecture 4: Interpretability and explainability in biomedical AI

HARVARD
MEDICAL SCHOOL

Marinka Zitnik
marinka@hms.harvard.edu

Responses to L3 Quick Check

Benefits of decentralized (federated) learning in healthcare?

1 - Generalizing to communities with smaller sample sizes
2 - Preserving the privacy of patients by protecting their critical identifying data
3 - Dealing with class imbalance in some populations

1. Helps maintain privacy of data by not requiring actual sharing of the data to build the models.
2. Good at handling biased local datasets by taking into account other parties' datasets.
3. Good at handling small sample sizes within individual local datasets.

- Federated learning reduces the cost of sending large amounts of data over a network, because only model parameters are transmitted rather than raw data.
- Federated learning improves model accuracy compared to local models, because the size of the training data set is increased.
- Federated learning provides greater protection for patient privacy than using pooled data, because patient data is not shared between different sites.

[1] Multiple institution can collaborate on data science projects by training models with their own data, without ever exchanging their own data or centralizing it.
[2] Protects the privacy of clients involved, in that, only model parameters are ever shared.
[3] The ability of pioneering research that is inclusive and incorporates diverse patient populations.

Responses to L3 Quick Check

Limitations of the study discussed in the class that used federated learning on EHRs for mortality prediction in hospitalized patients

The federated learning algorithm, as presented in the paper, is as slow as the slowest party. The aggregator must wait for all parties to return their updated weights. Some optimizations to the parallel processing could make the overall model train faster.

The final model is only as good as its data. The authors do not evaluate the model on a dataset from a different hospital system, which means the model might not generalize.

1. One limitation of the study regards its generalizability. The data used within federated learning and modeling was limited to Mount Sinai hospitals. Findings may then not be applicable to 'external' hospitals or other regions then.
2. A second limitation discussed was the data types used within the model. The model was limited to clinical data -- other data may have improved model performance and shed light on important features to consider.
3. A third limitation of the study is their lack of mention regarding the validity of the local datasets used. Federated learning inherently depends on the validity of local datasets, or how 'good' their data is. Poor validity may compromise overall model performance -- which the authors did not discuss.

Responses to L3 Quick Check

Limitations of the study discussed in the class that used federated learning on EHRs for mortality prediction in hospitalized patients

1. Federated learning may potentially be vulnerable to network-based cyber attacks compared to models trained in a purely offline setting
2. The authors did not explore the marginal benefit of federated learning for healthcare institutions with different amounts of local data. What is the threshold whereby models trained on local datasets only will perform equally well/better than federated learning models?

This is not specific to the federated learning model, but the authors focused on predicting outcomes over time using data from basically a single time point (admission). This means neglecting the wealth of data that accrues during the clinical course - accuracy of predictions at later time points could likely be improved by incorporating data on inpatient trajectory.

The authors in this study also benefited from their data coming from a single health system. Trying to harmonize measures across systems could be much more complicated - especially without sharing data centrally to notice that, for example, some hospitals label sodium "sodium" and some label it "Na," or that they have recorded height, weight or drip rates in different units. Not harmonizing data could result in some very erroneous model results. These challenges could be ameliorated with quality checks, but it would require human work and potentially some leakage of information between sites.

Outline for today's class

1. What is trustworthy ML and why should I care?

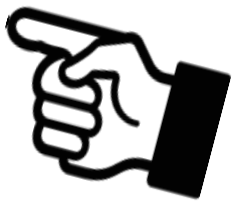

2. Interpretability vs. explainability
3. Explaining ML predictions
4. Case studies
   - Drug repurposing
   - Treatment recommendation

Trustworthy ML

- ML models are increasingly being deployed in real-world applications
  - It is critical to ensure that these models are behaving responsibly and are trustworthy
- There has been growing interest to develop and deploy ML models and algorithms that are:
  - Not only accurate
  - But also explainable, fair, privacy-preserving, causal, and robust
- This broad area of research is commonly referred to as trustworthy ML

Motivation

Model understanding is absolutely critical in several domains - particularly those involving **high stakes decisions**

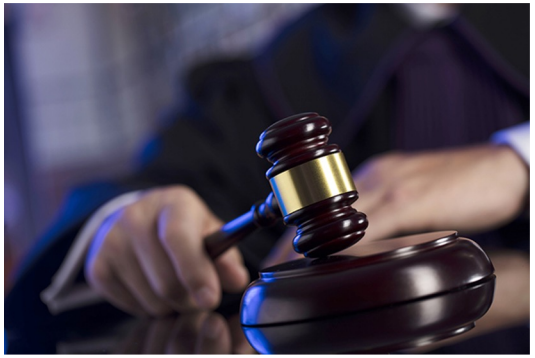

Why model understaning?

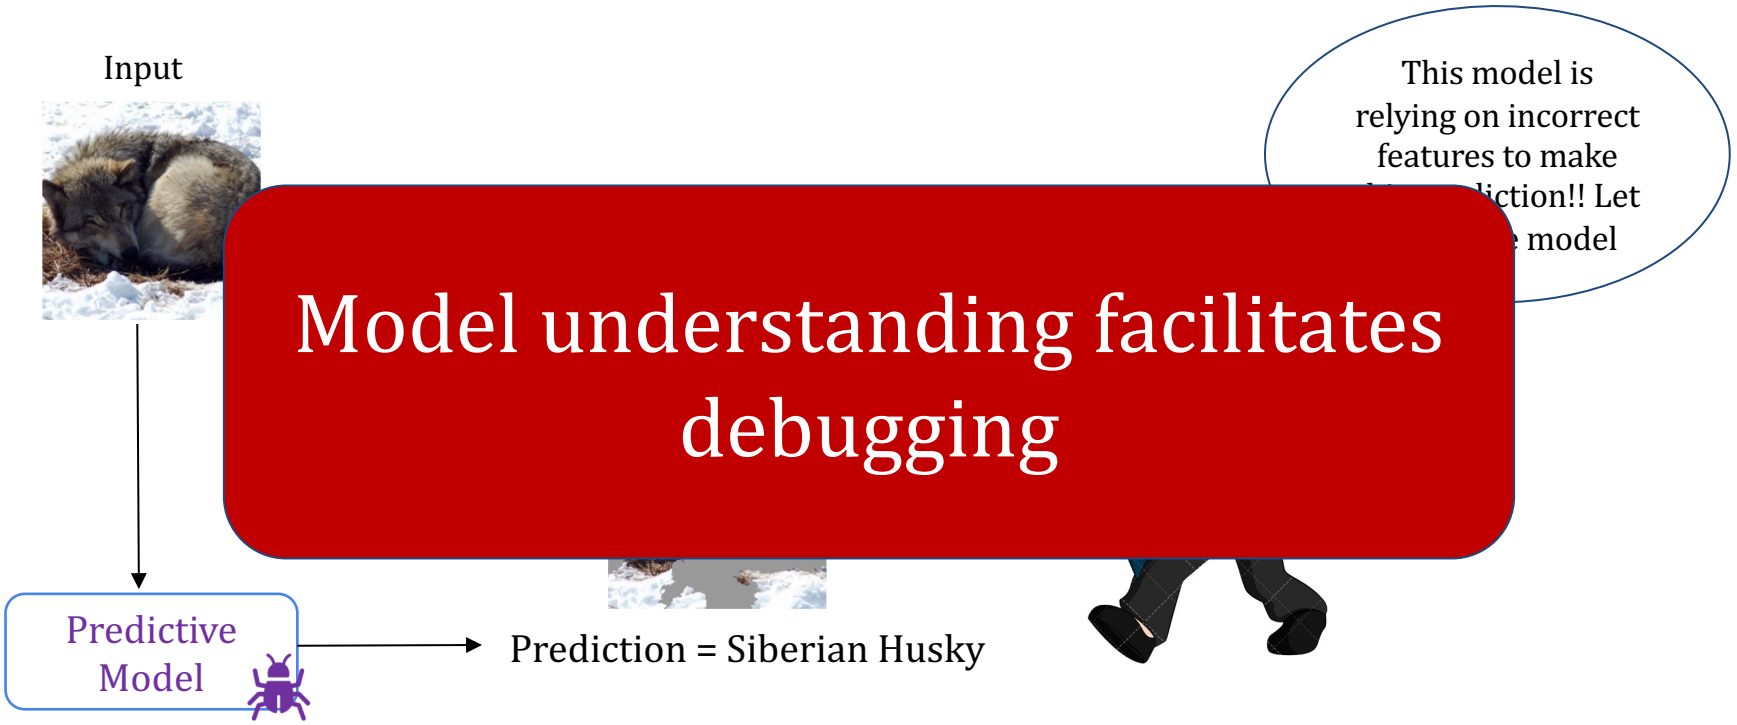

Why model understanding?

Why model understanding?

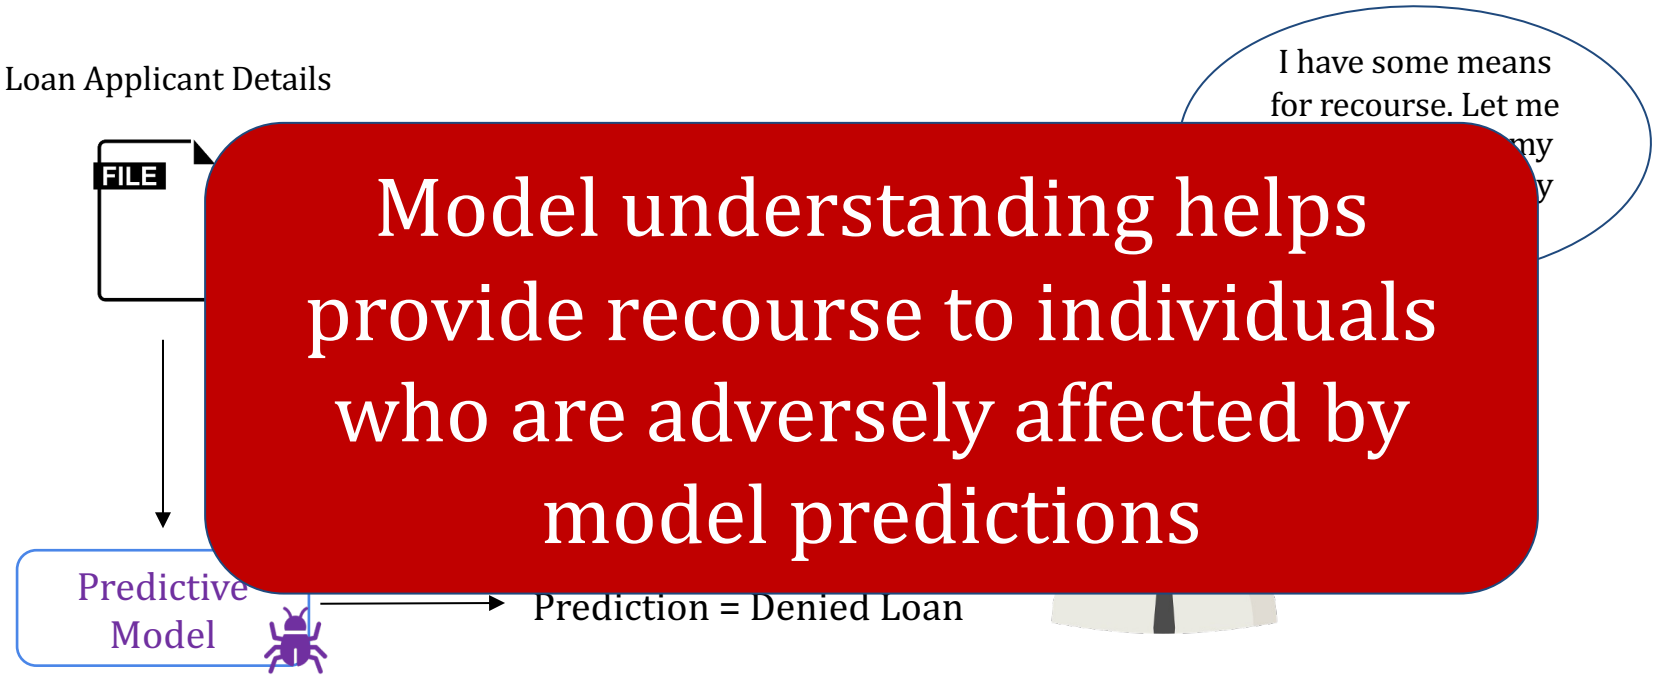

Model understanding helps provide recourse to individuals who are adversely affected by model predictions

Motivation: Why model understanding?

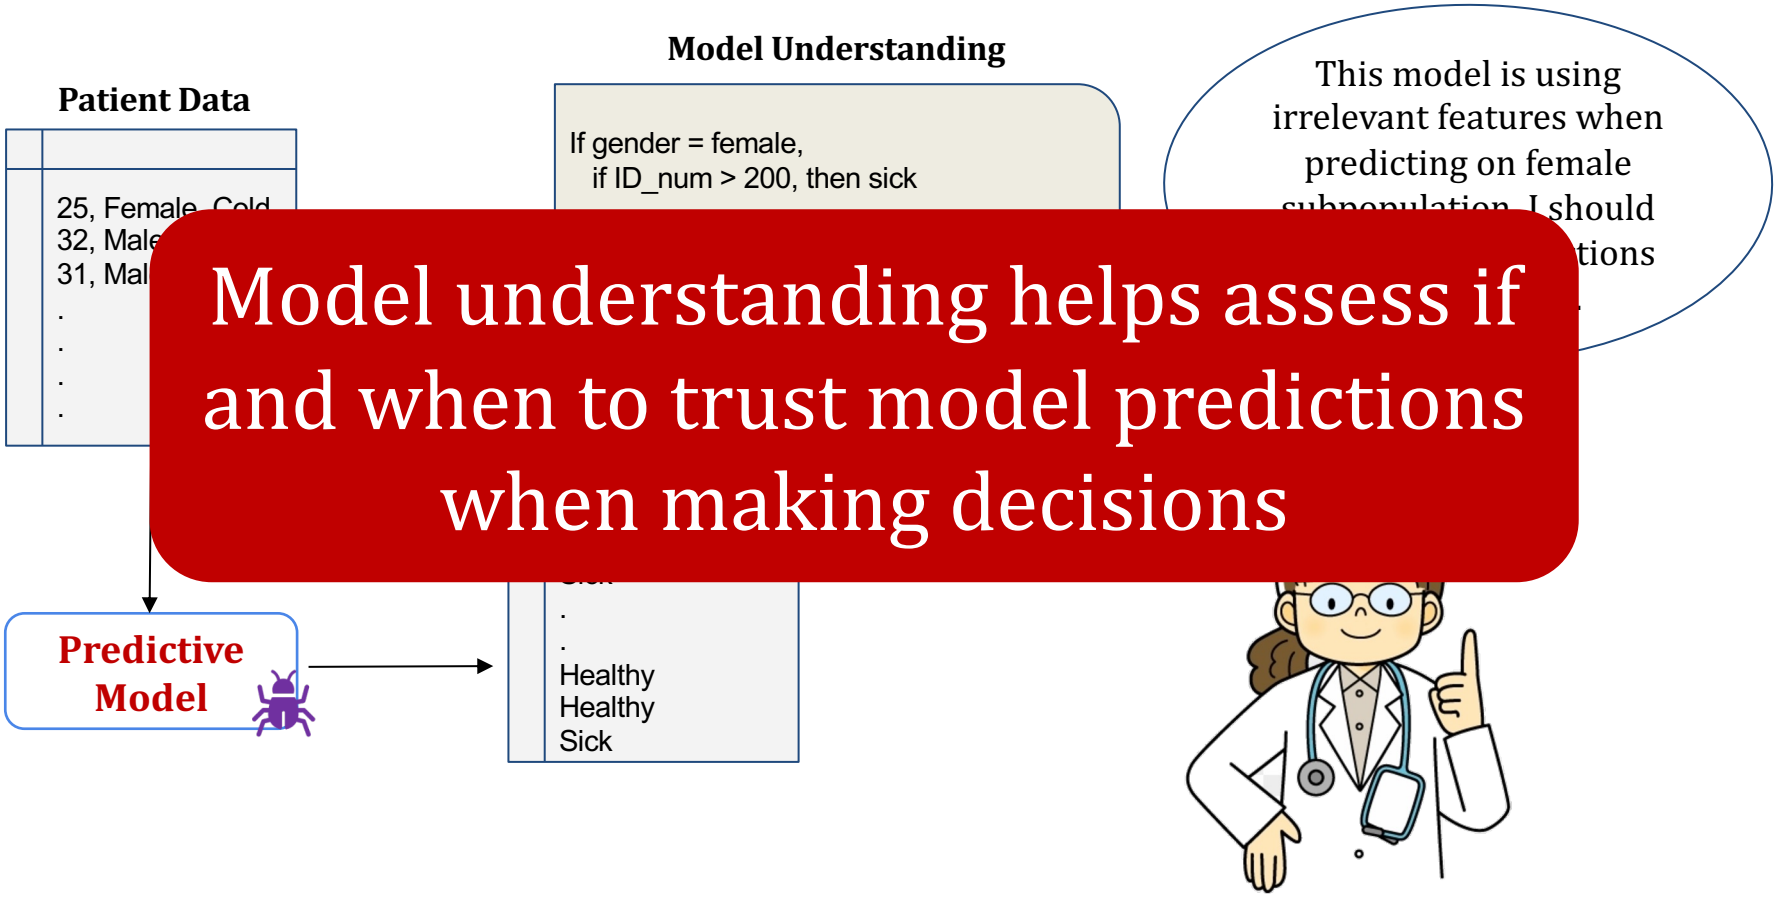

Motivation: Why model understanding?

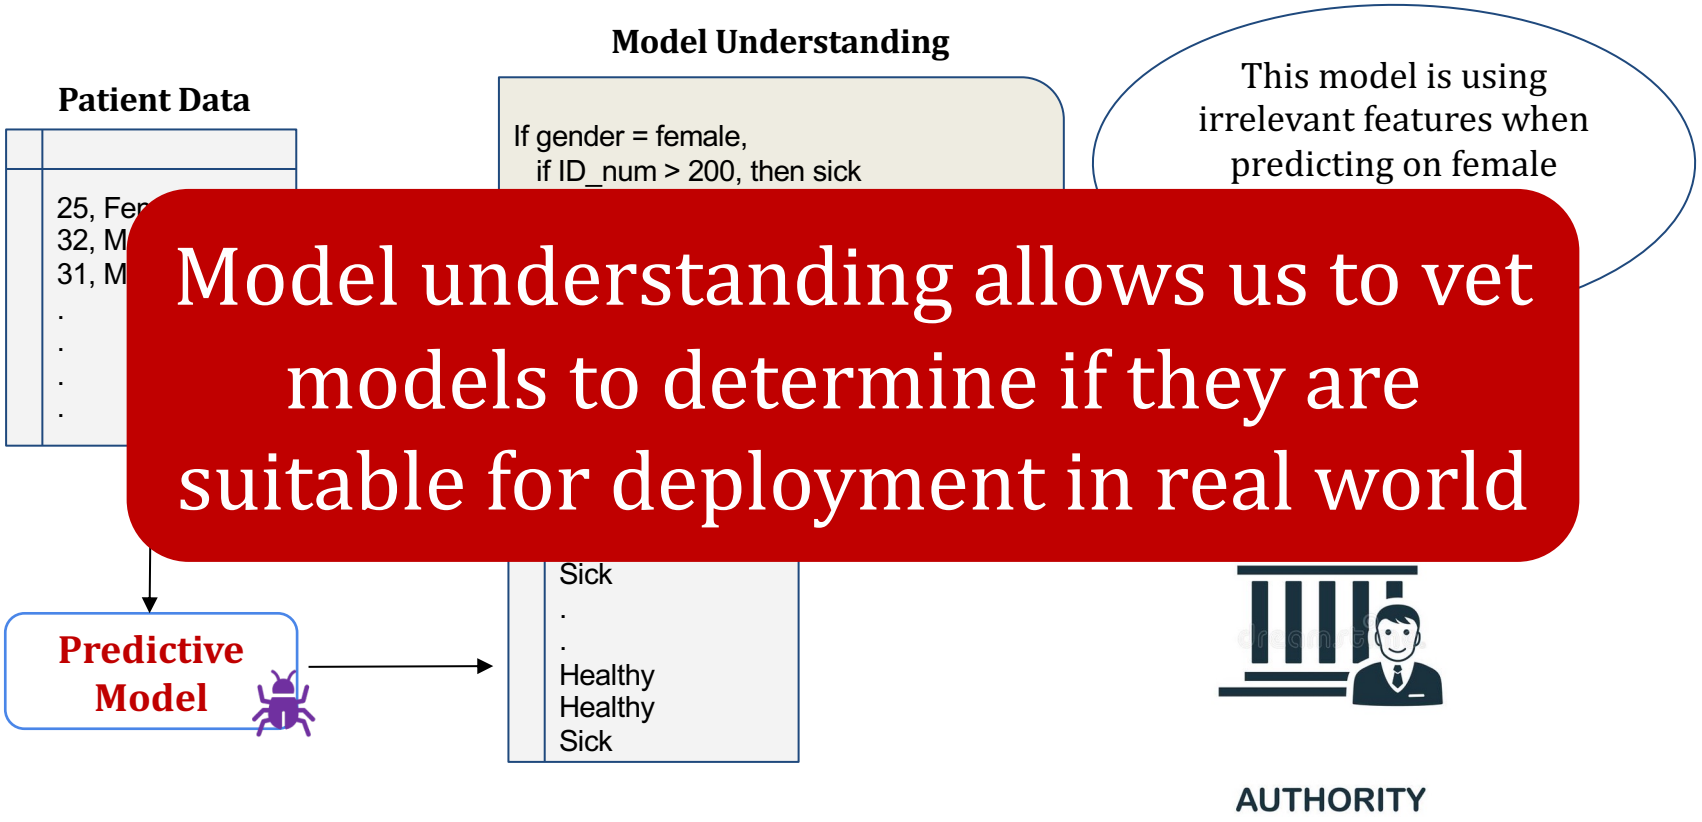

Why should I care about understanding ML models?

Utility

Debugging

Bias Detection

Recourse

If and when to trust model predictions

Vet models to assess suitability for deployment

Stakeholders

End users (e.g., loan applicants)

Decision makers (e.g., doctors, judges)

Regulatory agencies (e.g., FDA, European commission)

Researchers and engineers

Outline for today's class

1. What is trustworthy ML and why should I care?

2. Interpretability vs. explainability

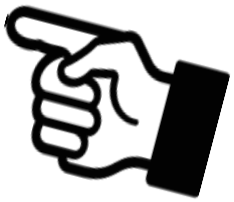

3. Explaining ML predictions

4. Case studies
   - Drug repurposing
   - Treatment recommendation

Achieving model understanding

Goal: Build inherently interpretable predictive models

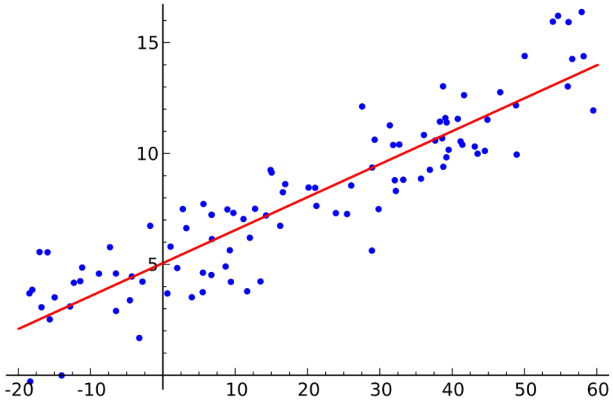

Linear regression

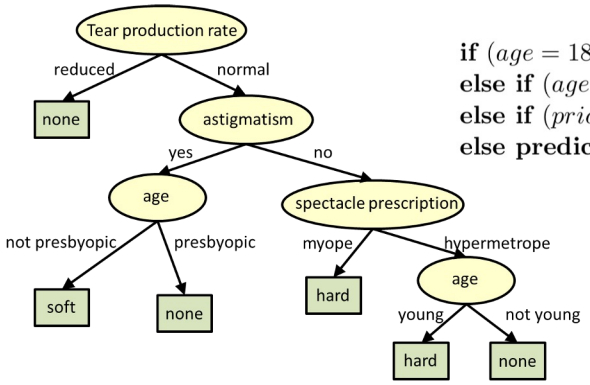

Decision trees

if $(age = 18 - 20)$ **and** $(sex = male)$ **then predict** yes
else if $(age = 21 - 23)$ **and** $(priors = 2 - 3)$ **then predict** yes
else if $(priors > 3)$ **then predict** yes
else predict no

Decision rules

Saliency map of a black box (deep learning) model does not explain anything except where the model is looking: We have no idea why this image is labeled as either a dog or a musical instrument when considering only saliency. The explanations look essentially the same for both classes

Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature Machine Intelligence* 2019

Inherently interpretable models vs. post hoc explanations

Accuracy-interpretability trade offs may exist in certain settings

Example

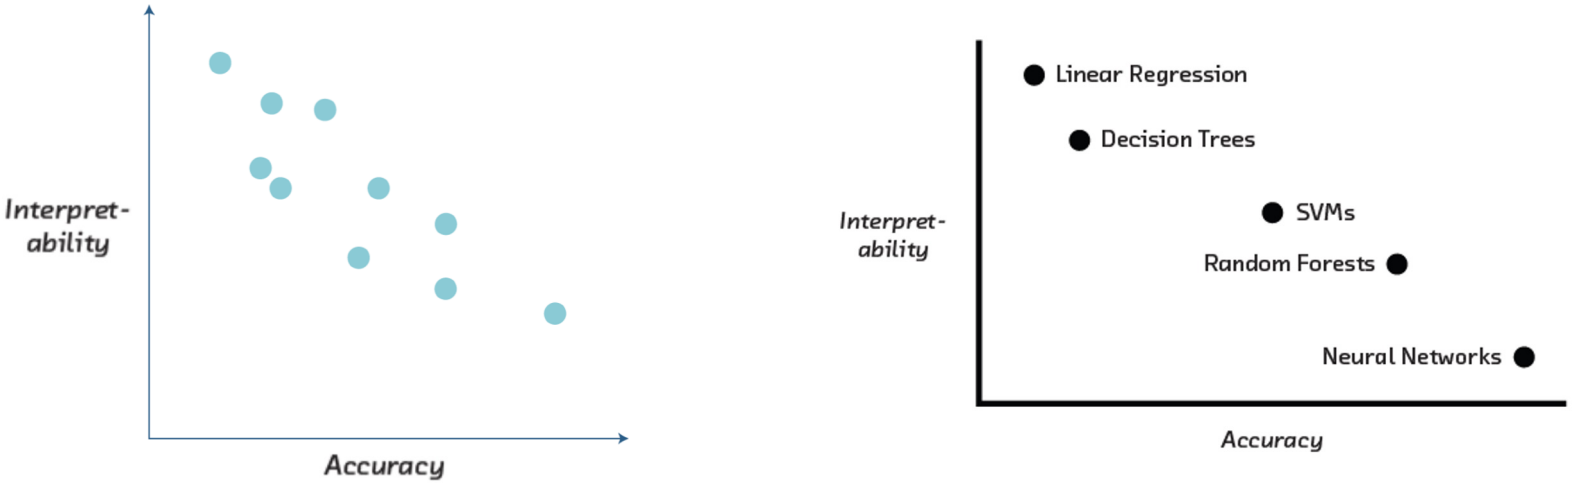

[Cireşan et. al. 2012, Caruana et. al. 2006, Frosst et. al. 2017, Stewart 2020]

Inherently interpretable models vs. post hoc explanations

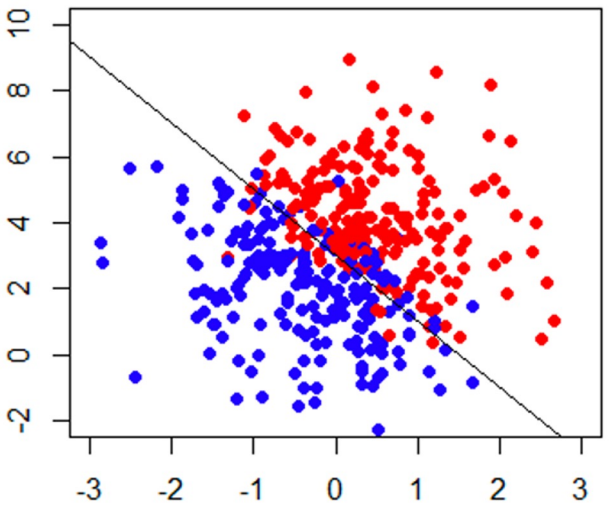

can build interpretable + accurate models

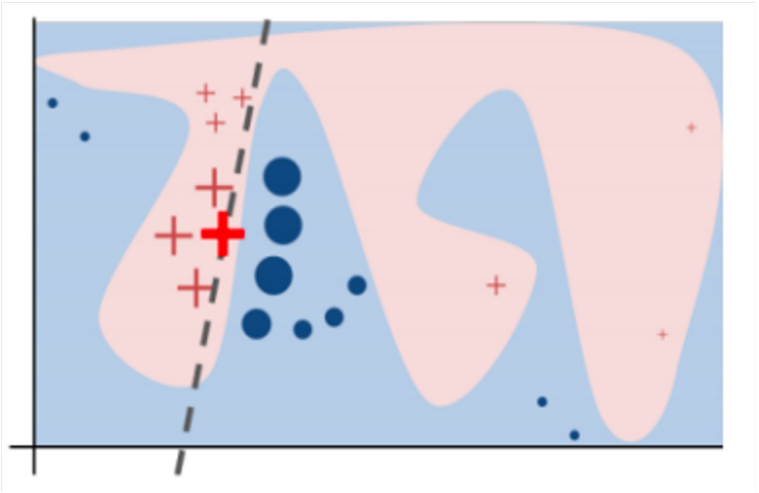

complex models might achieve higher accuracy

Achieving model understanding

Explain pre-built models *in a post-hoc manner*

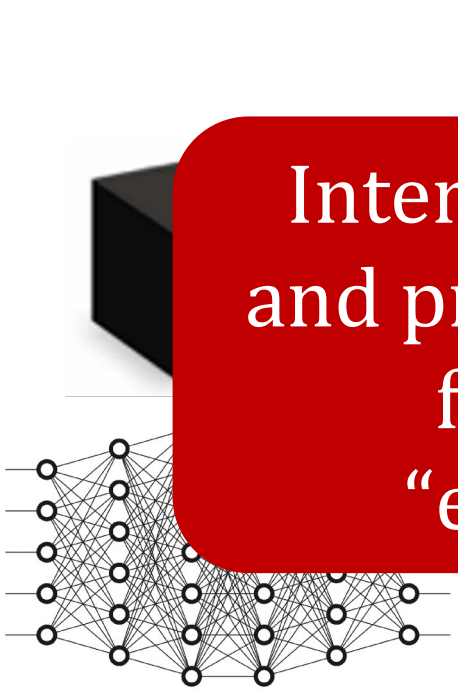

Interpretability/accuracy tradeoffs and proliferation of black box models force us to rely on post hoc "explanations" of ML models

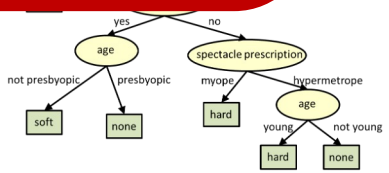

[Ribeiro et. al. 2016, 2018; Lakkaraju et. al. 2019]

Inherently interpretable models vs. post hoc explanations

- If you can build an interpretable model which is also adequately accurate for your setting, DO IT!

- Sometimes, you don't have enough data to build your model from scratch
- And, all you have is a (proprietary) black box!
- Post hoc explanations come to the rescue!

Next: Overview of post hoc explanations methods

Outline for today's class

1. What is trustworthy ML and why should I care?
2. Interpretability vs. explainability
3. Explaining ML predictions
4. Case studies
   - Drug repurposing
   - Treatment recommendation

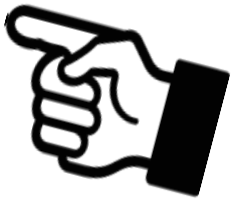

Explainable AI

"Explainable AI refers to the set of approaches that provide an interpretable description of the behavior of a given (complex) model to end users."

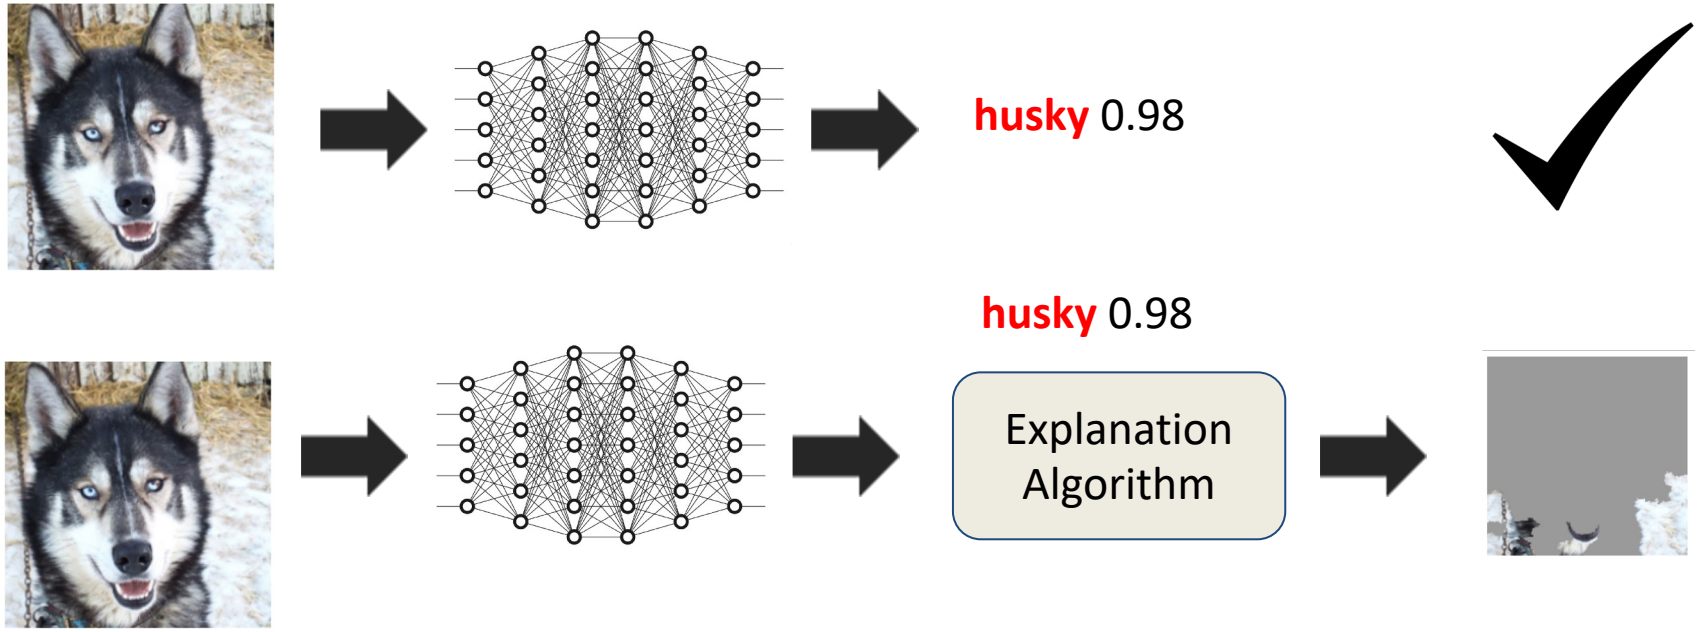

What is an explanation?

- **Definition:** Interpretable description of the model behavior

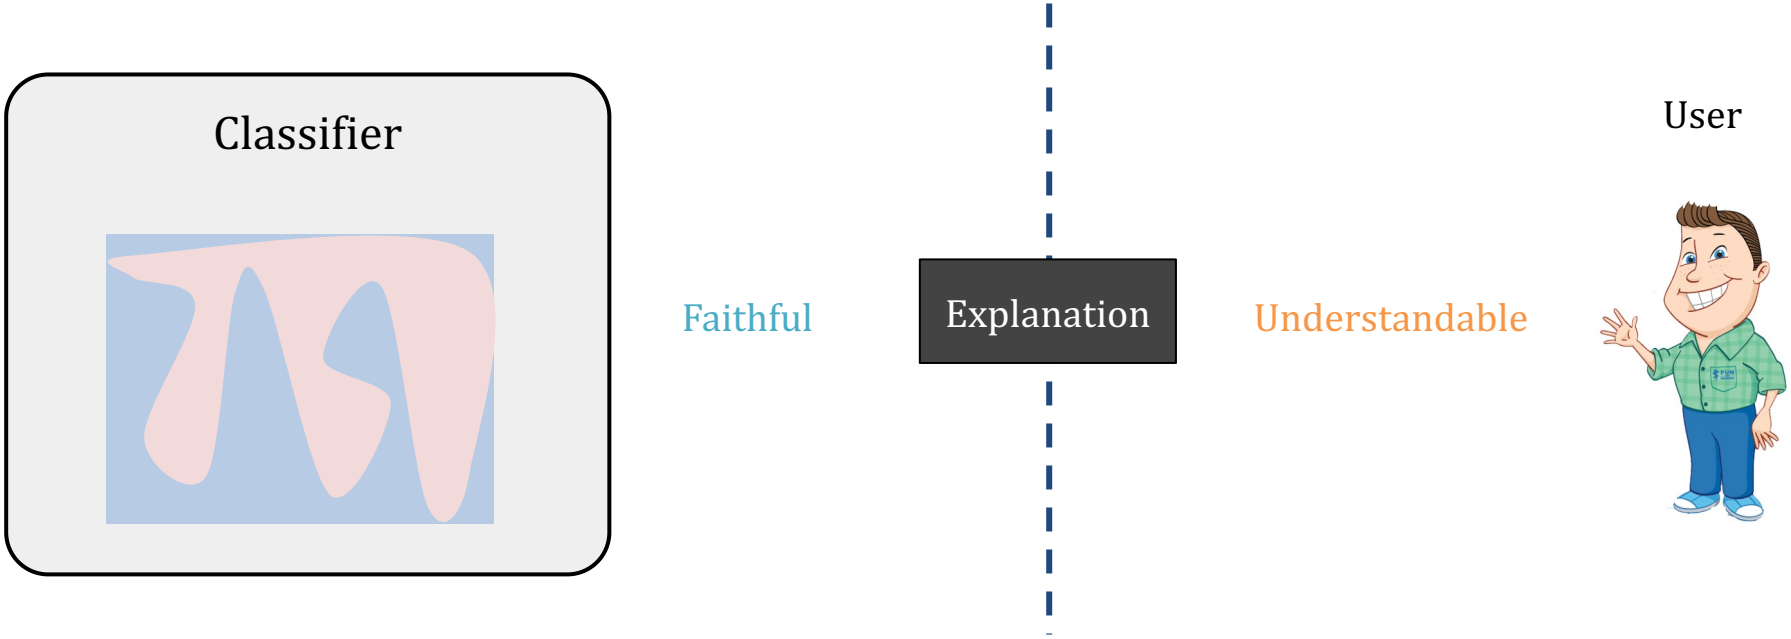

Overview of explanation methods

Local explanations

Explain individual predictions

Help unearth biases in the *local neighborhood* of a given instance

Help vet if individual predictions are being made for the right reasons

Global explanations

Explain complete behavior of the model

Sheds light on *big picture biases* affecting larger subgroups

Help vet if the model, at a high level, is suitable for deployment

Overview of explanation methods

- Local explanation methods:
  - Feature importance scoring
  - Integrated gradients
  - Prototype explanations
  - Counterfactuals

- Global explanation methods:
  - Collection of local explanations
  - Representation-based explanations
  - Model distillation

LIME: Local interpretable model-agnostic explanations

1. Sample points around x_i

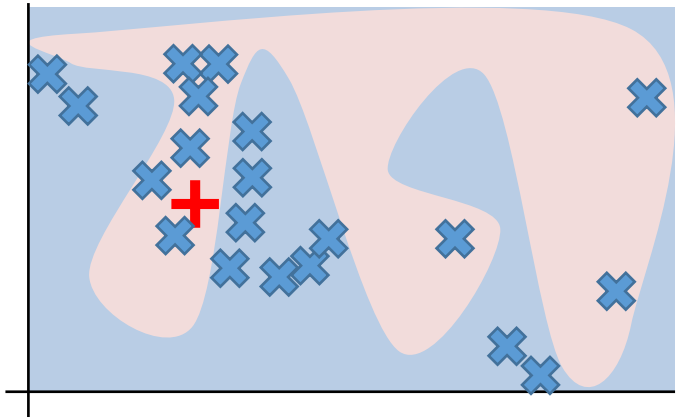

[Ribeiro et al. 2016]

LIME: Local interpretable model-agnostic explanations

1. Sample points around x_i
2. Use model to predict labels for each sample

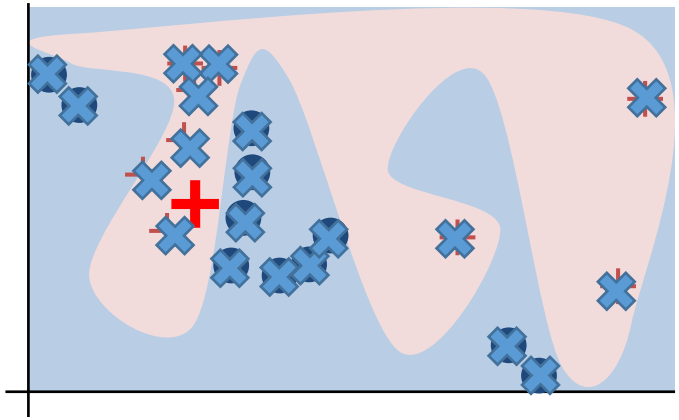

[Ribeiro et al. 2016]

LIME: Local interpretable model-agnostic explanations

1. Sample points around x_i
2. Use model to predict labels for each sample
3. Weigh samples according to distance to x_i

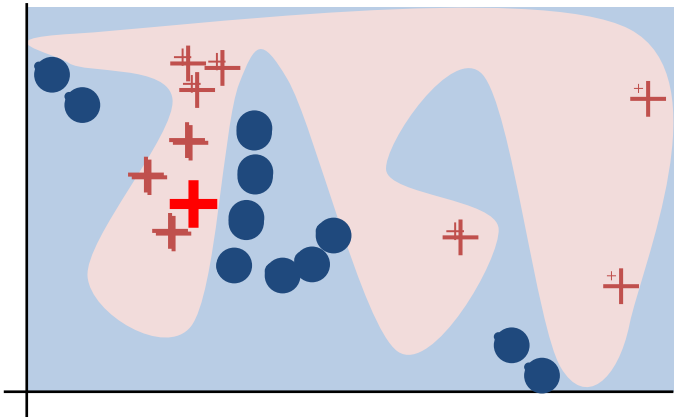

[Ribeiro et al. 2016]

LIME: Local interpretable model-agnostic explanations

1. Sample points around x_i
2. Use model to predict labels for each sample
3. Weigh samples according to distance to x_i
4. Learn simple linear model on weighted samples

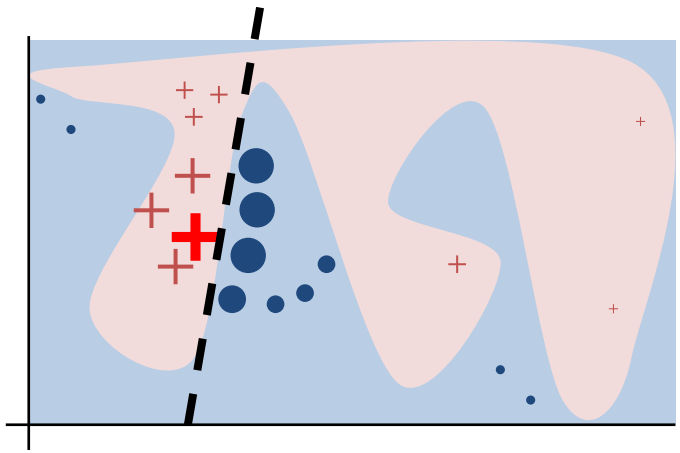

[Ribeiro et al. 2016]

LIME: Local interpretable model-agnostic explanations

1. Sample points around x_i
2. Use model to predict labels for each sample
3. Weigh samples according to distance to x_i
4. Learn simple linear model on weighted samples
5. Use simple linear model to explain x_i

Another popular method which outputs feature importance scores: SHAP

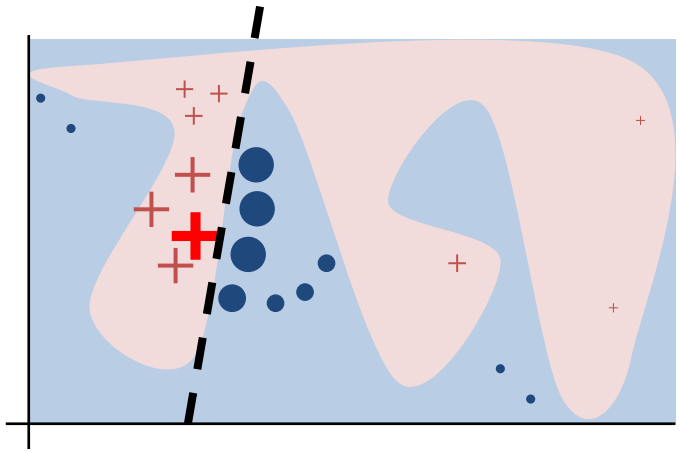

[Ribeiro et al. 2016]

Overview of explanation methods

- **Local explanation methods:**
  - Feature importance scoring
  - Integrated gradients
  - Prototype explanations
  - Counterfactuals

- **Global explanation methods:**
  - Collection of local explanations
  - Representation-based explanations
  - Model distillation

Integrated Gradients (IG)

- Integrated Gradients (IG) is an explanation method for deep neural networks
- It identifies important features that contribute most to the model's prediction

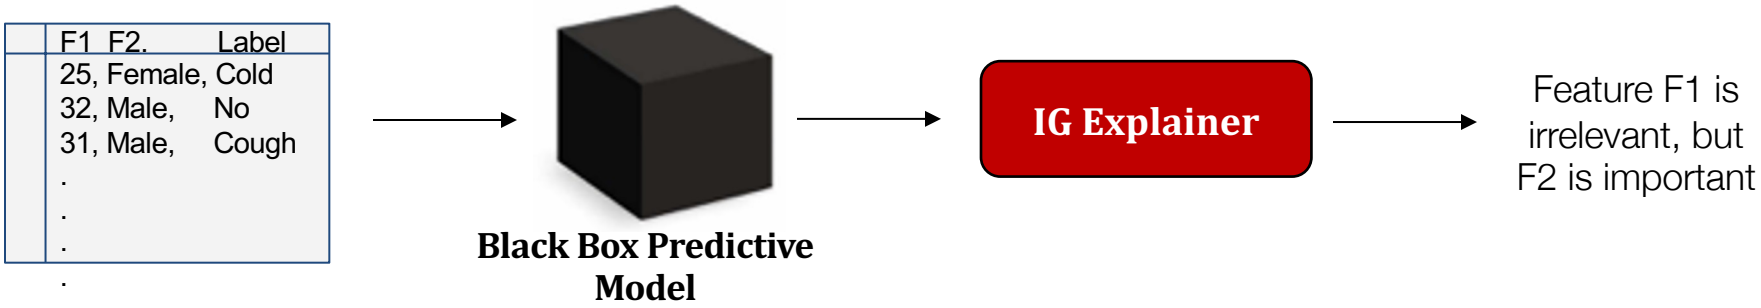

- Appealing properties of integrated gradients:
  - It can be applied to any differentiable model like models for images, text, or structured data
  - It requires no modification to the original ML model

Axiomatic Attribution for Deep Networks, *ICML 2017*

https://distill.pub/2020/attribution-baselines

How does IG work?

- IG computes gradients of the model's prediction w.r.t. input features
- IG is built on two axioms which need to be satisfied:
  - Sensitivity and
  - Implementation invariance
- Sensitivity:
  - We establish a baseline instance as a starting point
  - We then build a sequence of instances which we interpolate from a baseline instance to the actual instance to calculate
- Implementation invariance:
  - Implementation invariance is satisfied when two functionally equivalent models have identical attributions for the same input image and the baseline image.
  - Two models are functionally equivalent when their outputs are equal for all inputs despite having very different implementations

Axiomatic Attribution for Deep Networks, *ICML 2017*

Calculating and visualizing IG

- **Setup:**
  - Let's consider an ML model for image classification
  - We aim to use IG to explain the predicted image label

- **Step 1:**
  - Start from the baseline where baseline can be a black image whose pixel values are all zero or an all-white image, or a random image
  - Baseline input is one where the prediction is neutral and is central to any explanation method and visualizing pixel feature importances

Calculating and visualizing IG

- Step 2:
  - Generate a linear interpolation between the baseline and the original image
  - Interpolated images are small steps(α) in the feature space between your baseline and input image and consistently increase with each interpolated image's intensity

Axiomatic Attribution for Deep Networks, *ICML 2017*

Calculating and visualizing IG

- Step 3: Calculate gradients to measure the relationship between changes to a feature and changes in the model's predictions
- The gradient informs which pixel has the strongest effect on the models predicted class probabilities
  - Varying variable changes the output, and the variable will receive some attribution to help calculate the feature importances for the input image
  - Variable that does not affect the output gets no attribution
- Step 4: Compute the numerical approximation through averaging gradients (that's why the method's name is integrated gradients)

Axiomatic Attribution for Deep Networks, *ICML 2017*

Calculating and visualizing IG

- Step 5:
  - Scale IG to the input image to ensure that the attribution values are accumulated across multiple interpolated images are all in the same units
  - Represent the IG on the input image with the pixel importances

IG helps us explain what an ML model looks at to make a prediction by highlighting the feature importances.
It does this by computing the gradient of the model's prediction output to its input features.

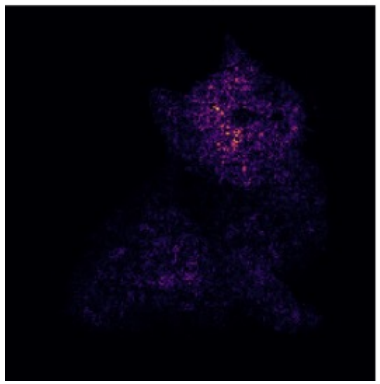

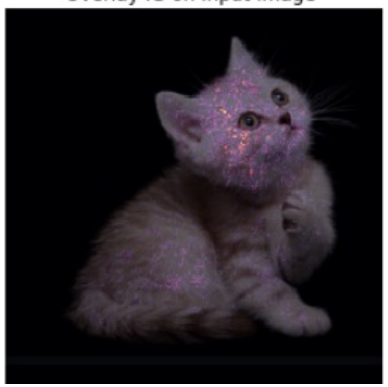

Axiomatic Attribution for Deep Networks, *ICML 2017*

Overview of explanation methods

- **Local explanation methods:**
  - Feature importance scoring
  - Integrated gradients
  - **Prototype explanations**
  - Counterfactuals

- **Global explanation methods:**
  - Collection of local explanations
  - Representation-based explanations
  - Model distillation

Prototype-based explanations

- Use examples (synthetic or natural) to explain individual predictions

- **Influence Functions (Koh & Liang 2017)**
  - Identify instances in the training set that are responsible for the prediction of a given test instance

- **Activation Maximization (Erhan et al. 2009)**
  - Identify examples (synthetic or natural) that strongly activate a function (neuron) of interest

Overview of explanation methods

- **Local explanation methods:**
  - Feature importance scoring
  - Integrated gradients
  - Prototype explanations
  - Counterfactuals

- **Global explanation methods:**
  - Collection of local explanations
  - Representation-based explanations
  - Model distillation

Counterfactual explanations

What features need to be changed and by how much to flip a model's prediction?

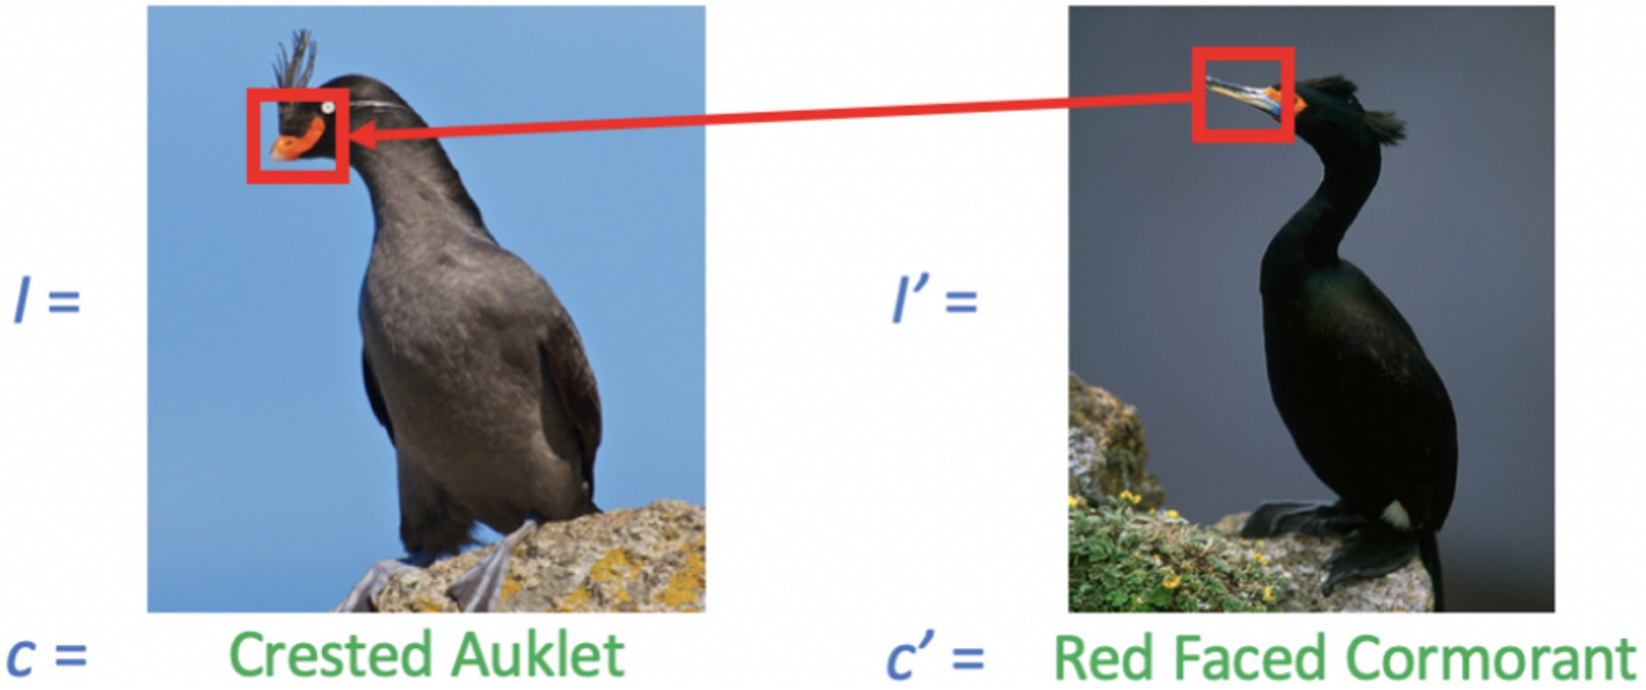

[Goyal et. al., 2019]

Counterfactual explanations

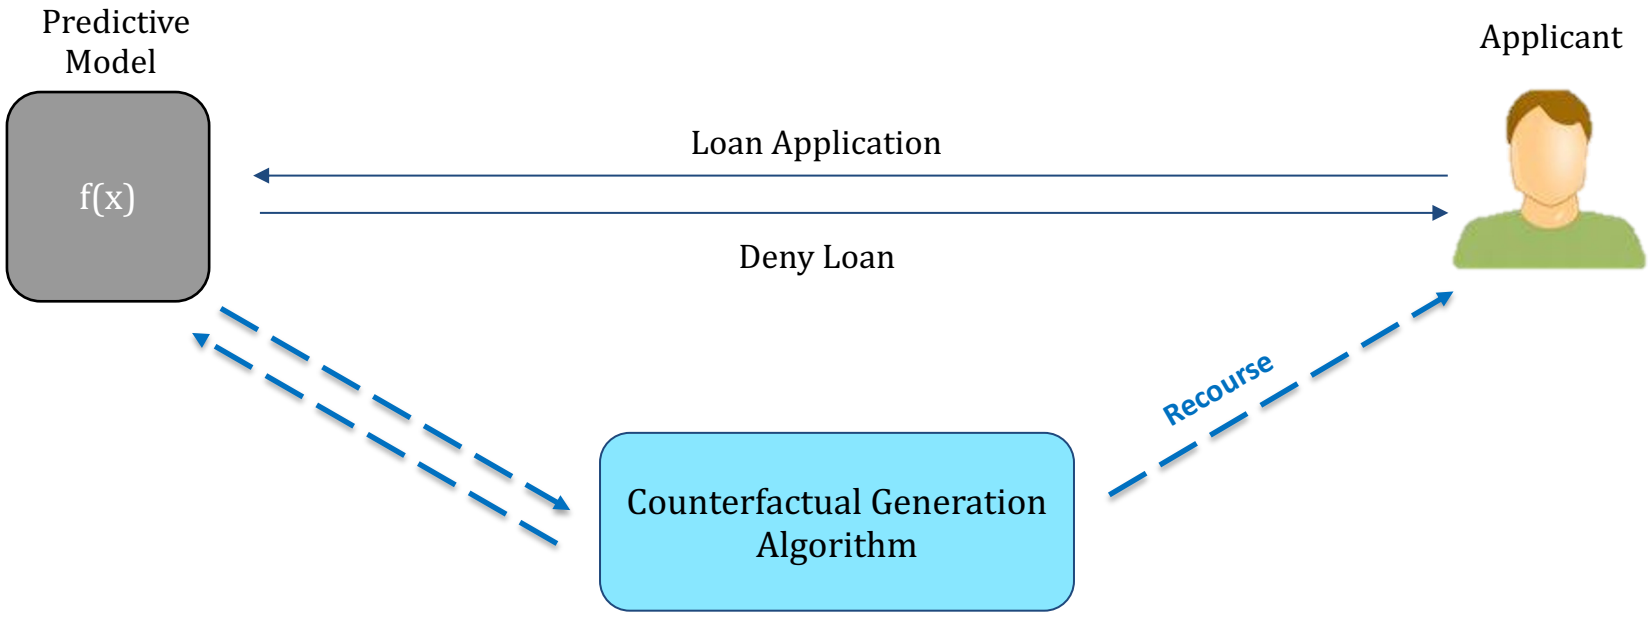

Recourse: Increase your salary by 50K & pay your credit card bills on time for next 3 months

Generating counterfactual explanations: Intuition

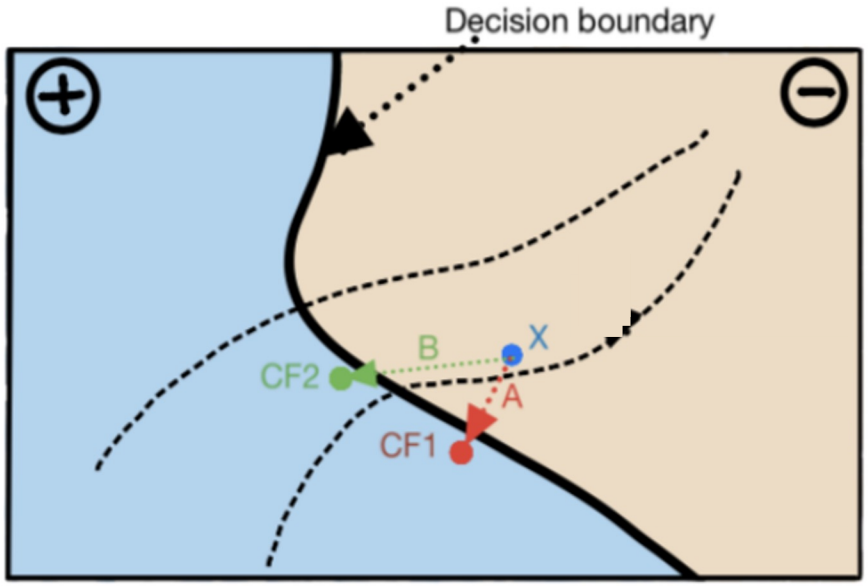

Proposed solutions differ on:

1. How to choose among candidate counterfactuals?

1. How much access is needed to the underlying predictive model?

[Verma et. al., 2020]

Quick Check

https://forms.gle/An2ZzQHbc568XAhe9

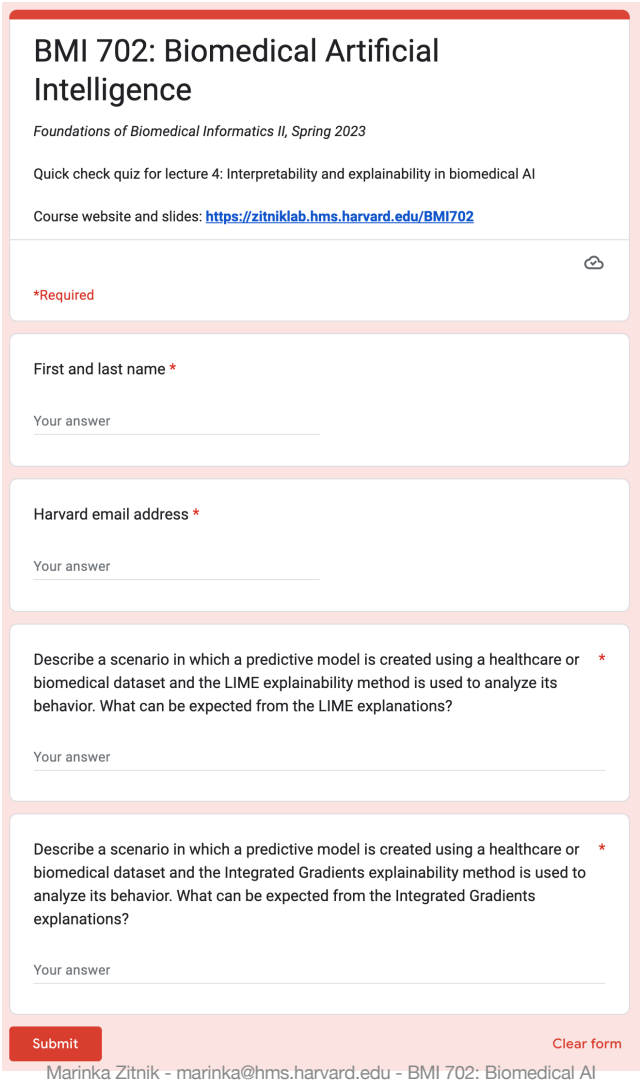

BMI 702: Biomedical Artificial Intelligence

Foundations of Biomedical Informatics II, Spring 2023

Quick check quiz for lecture 4: Interpretability and explainability in biomedical AI

Course website and slides: https://zitniklab.hms.harvard.edu/BMI702

*Required

First and last name *

Your answer

Harvard email address *

Your answer

Describe a scenario in which a predictive model is created using a healthcare or biomedical dataset and the LIME explainability method is used to analyze its behavior. What can be expected from the LIME explanations? *

Your answer

Describe a scenario in which a predictive model is created using a healthcare or biomedical dataset and the Integrated Gradients explainability method is used to analyze its behavior. What can be expected from the Integrated Gradients explanations? *

Your answer

Submit

Clear form

Overview of explanation methods

- Local explanation methods:
  - Feature importance scoring
  - Integrated gradients
  - Prototype explanations
  - Counterfactuals

- Global explanation methods:
  - Collection of local explanations
  - Representation-based explanations
  - Model distillation

Global explanations from local feature importances: SP-LIME

LIME explains a single prediction
local behavior for a single instance

Can't examine all explanations
Instead pick *k* explanations to show to the user

Representative
Should summarize the model's global behavior

Diverse
Should not be redundant in their descriptions

SP-LIME uses submodular optimization and *greedily* picks k explanations

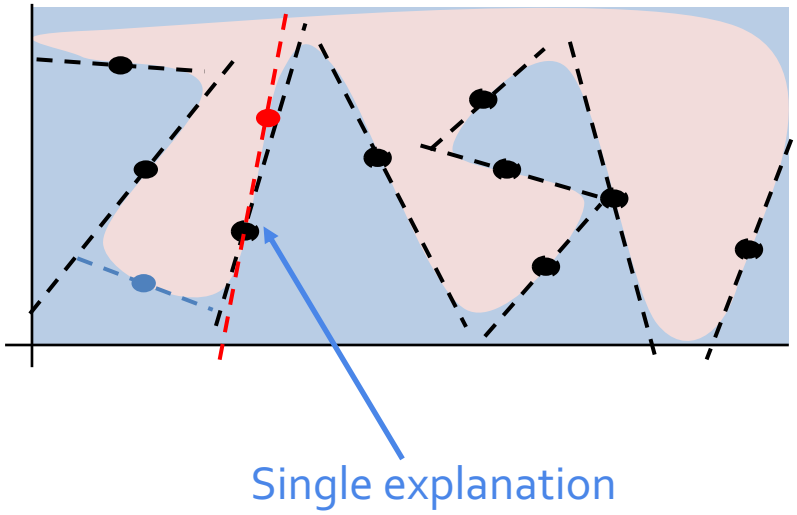

[Ribeiro et. al., 2016]

Overview of explanation methods

- Local explanation methods:
  - Feature importance scoring
  - Integrated gradients
  - Prototype explanations
  - Counterfactuals

- Global explanation methods:
  - Collection of local explanations
  - Representation-based explanations
  - Model distillation

Representation-based explanations

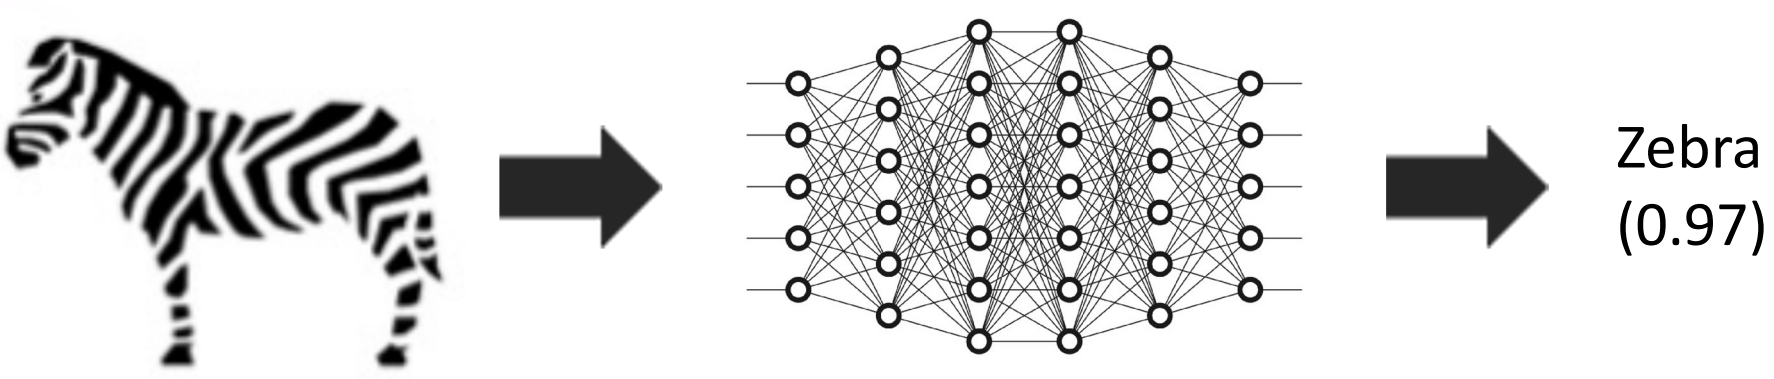

How important is the notion of “stripes” for this prediction?

[Kim et. al., 2018]

Representation-based explanations: TCAV approach

Examples of the concept “stripes”

Random examples

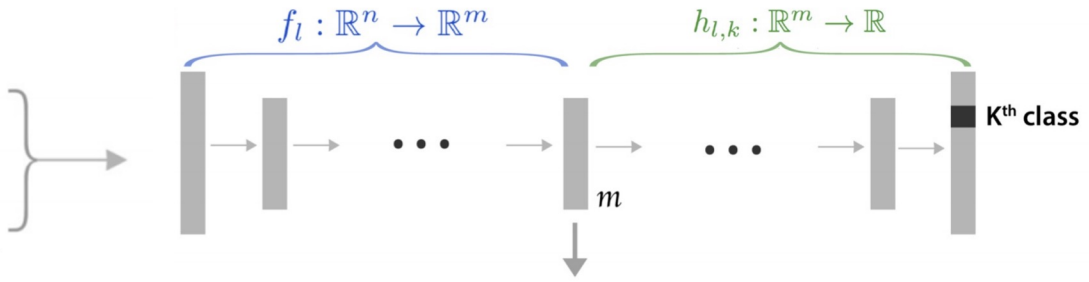

Train a linear classifier to separate activations

The vector orthogonal to the decision boundary denotes the concept “stripes”

Compute gradient w.r.t. this vector to determine how important is the notion of stripes for a prediction

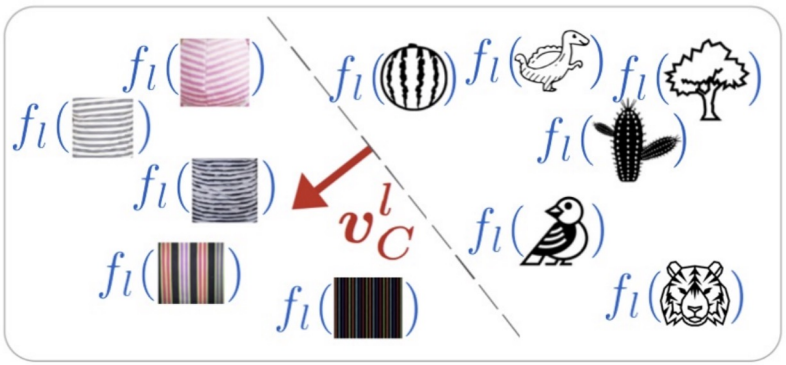

[Kim et. al., 2018]

TCAV = testing with concept activation vectors

Overview of explanation methods

- **Local explanation methods:**
  - Feature importance scoring
  - Integrated gradients
  - Prototypes/Example-based explanations
  - Counterfactuals

- **Global explanation methods:**
  - Collection of local explanations
  - Representation-based explanations
  - **Model distillation**

Model distillation

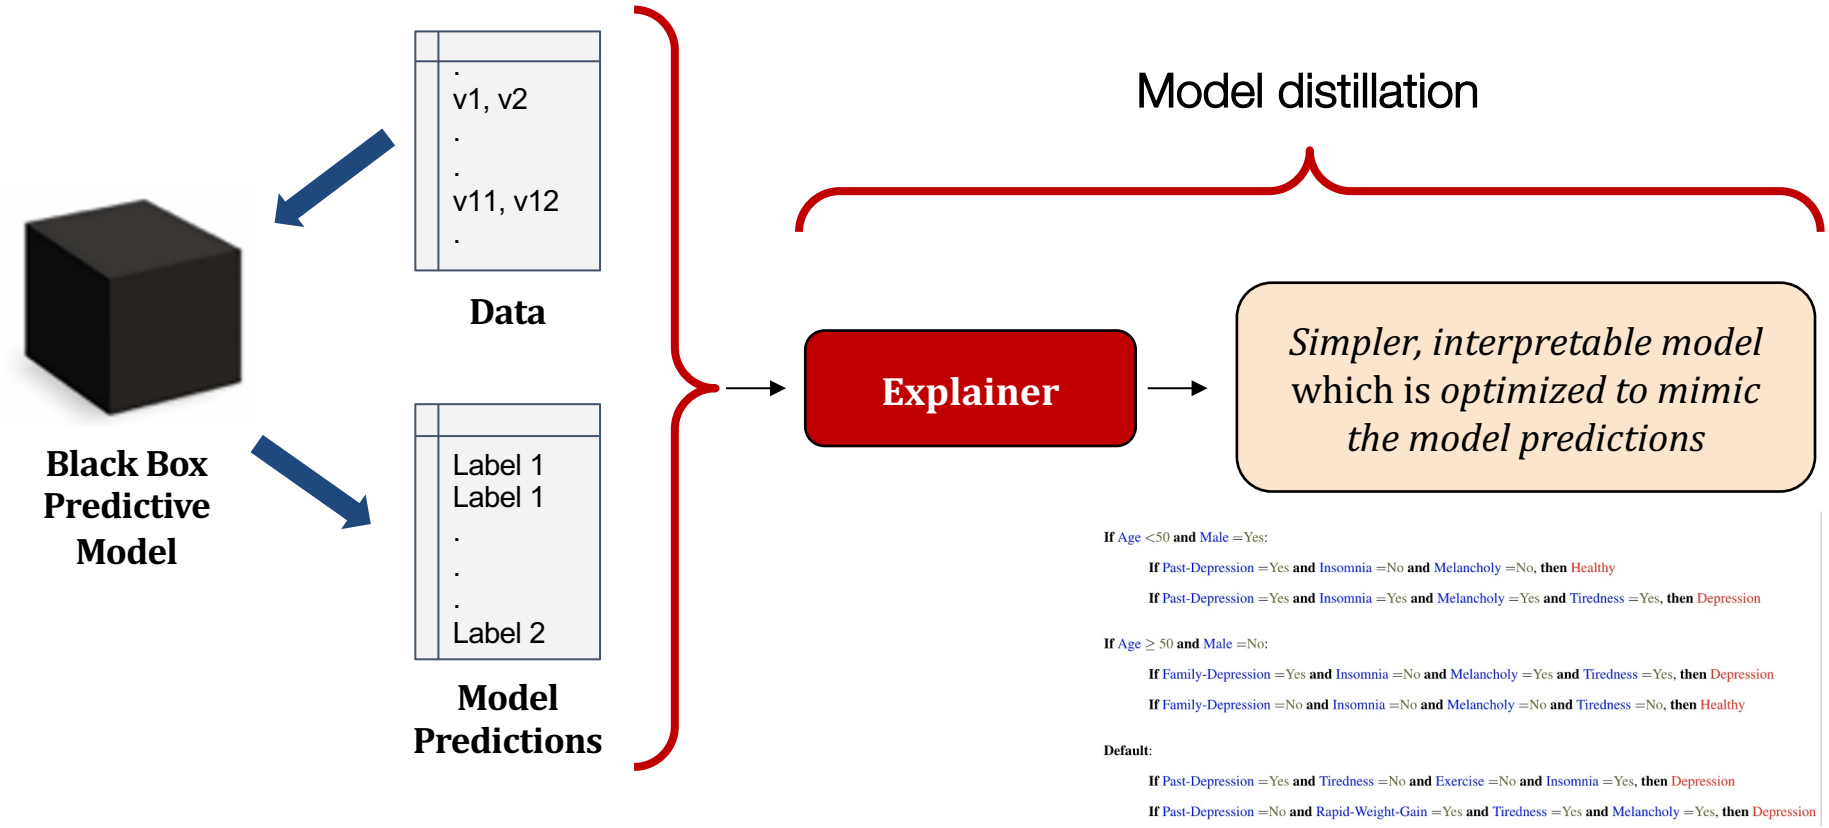

Model distillation using decision trees

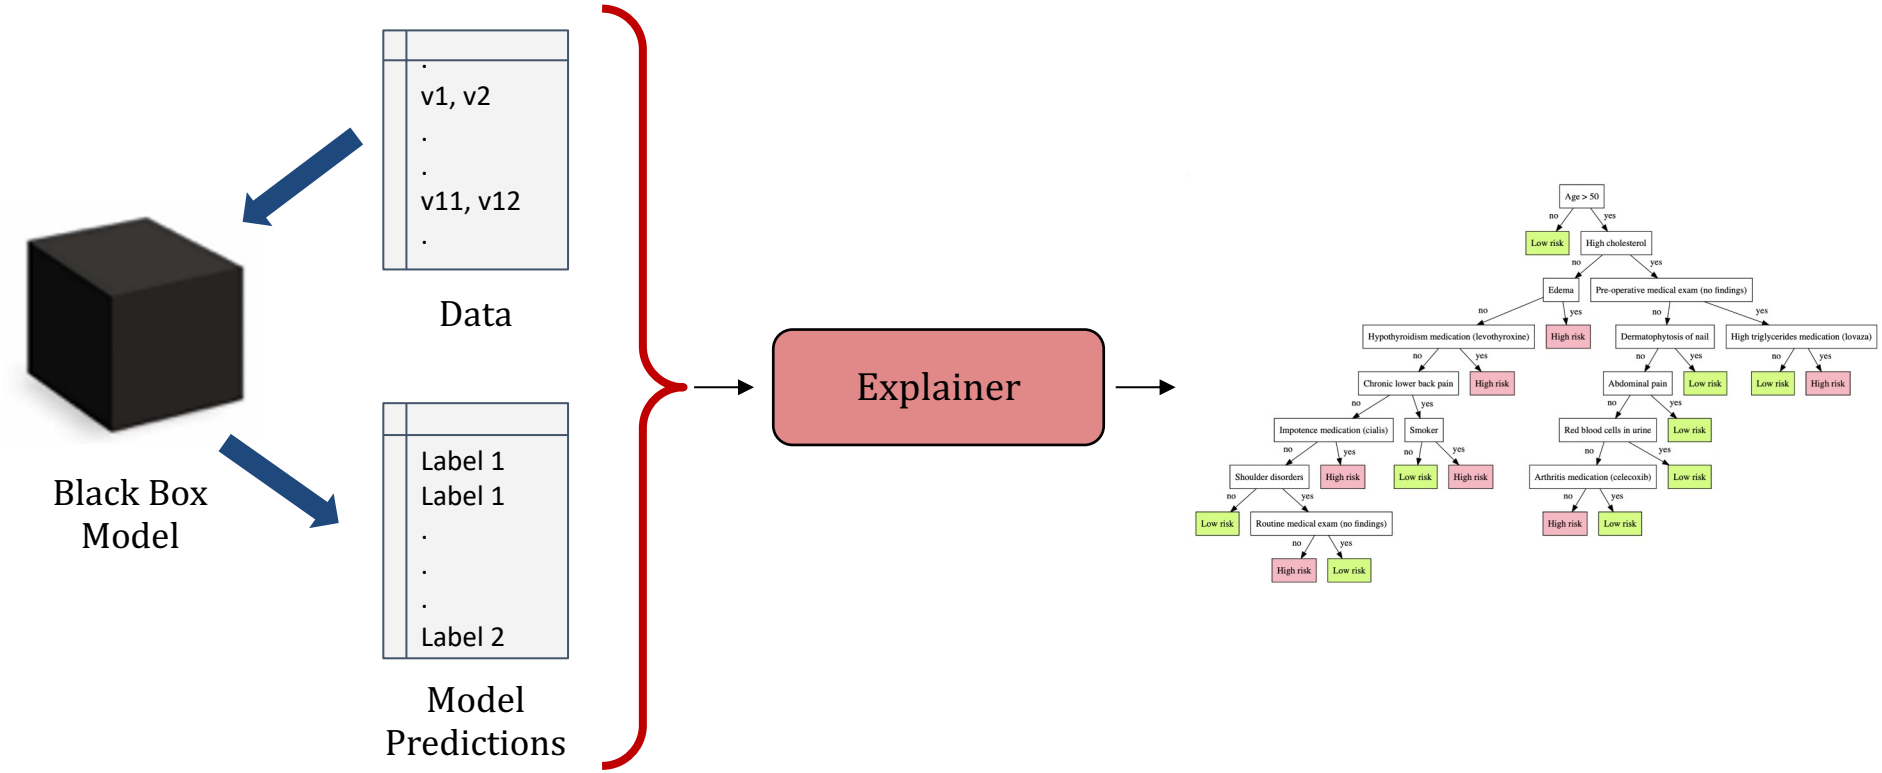

[Bastani et. al., 2019]

Model distillation using decision sets

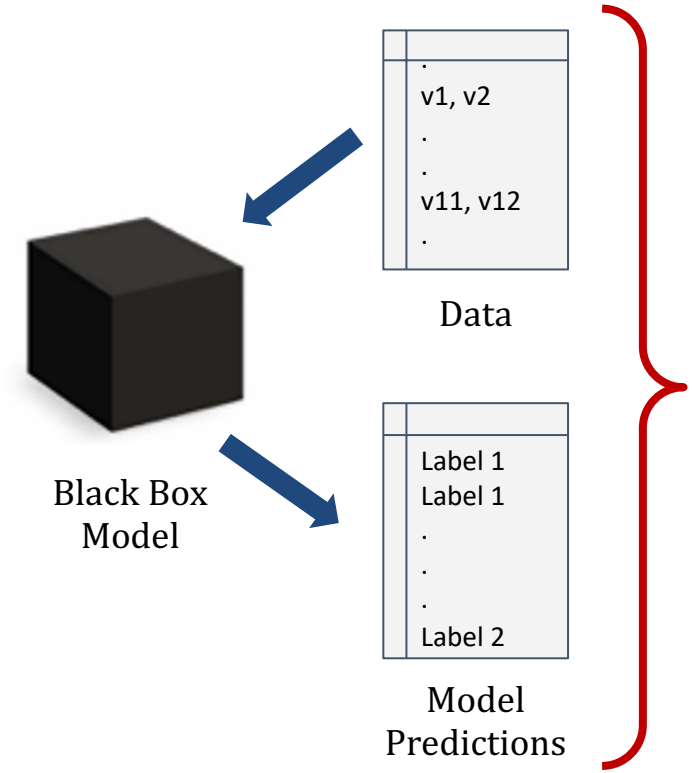

Explainer

If Age <50 **and** Male =Yes:

If Past-Depression =Yes **and** Insomnia =No **and** Melancholy =No, **then** Healthy

If Past-Depression =Yes **and** Insomnia =Yes **and** Melancholy =Yes **and** Tiredness =Yes, **then** Depression

If Age ≥ 50 **and** Male =No:

If Family-Depression =Yes **and** Insomnia =No **and** Melancholy =Yes **and** Tiredness =Yes, **then** Depression

If Family-Depression =No **and** Insomnia =No **and** Melancholy =No **and** Tiredness =No, **then** Healthy

Default:

If Past-Depression =Yes **and** Tiredness =No **and** Exercise =No **and** Insomnia =Yes, **then** Depression

If Past-Depression =No **and** Rapid-Weight-Gain =Yes **and** Tiredness =Yes **and** Melancholy =Yes, **then** Depression

[Lakkaraju et. al., 2019]

Model distillation using generalized additive models

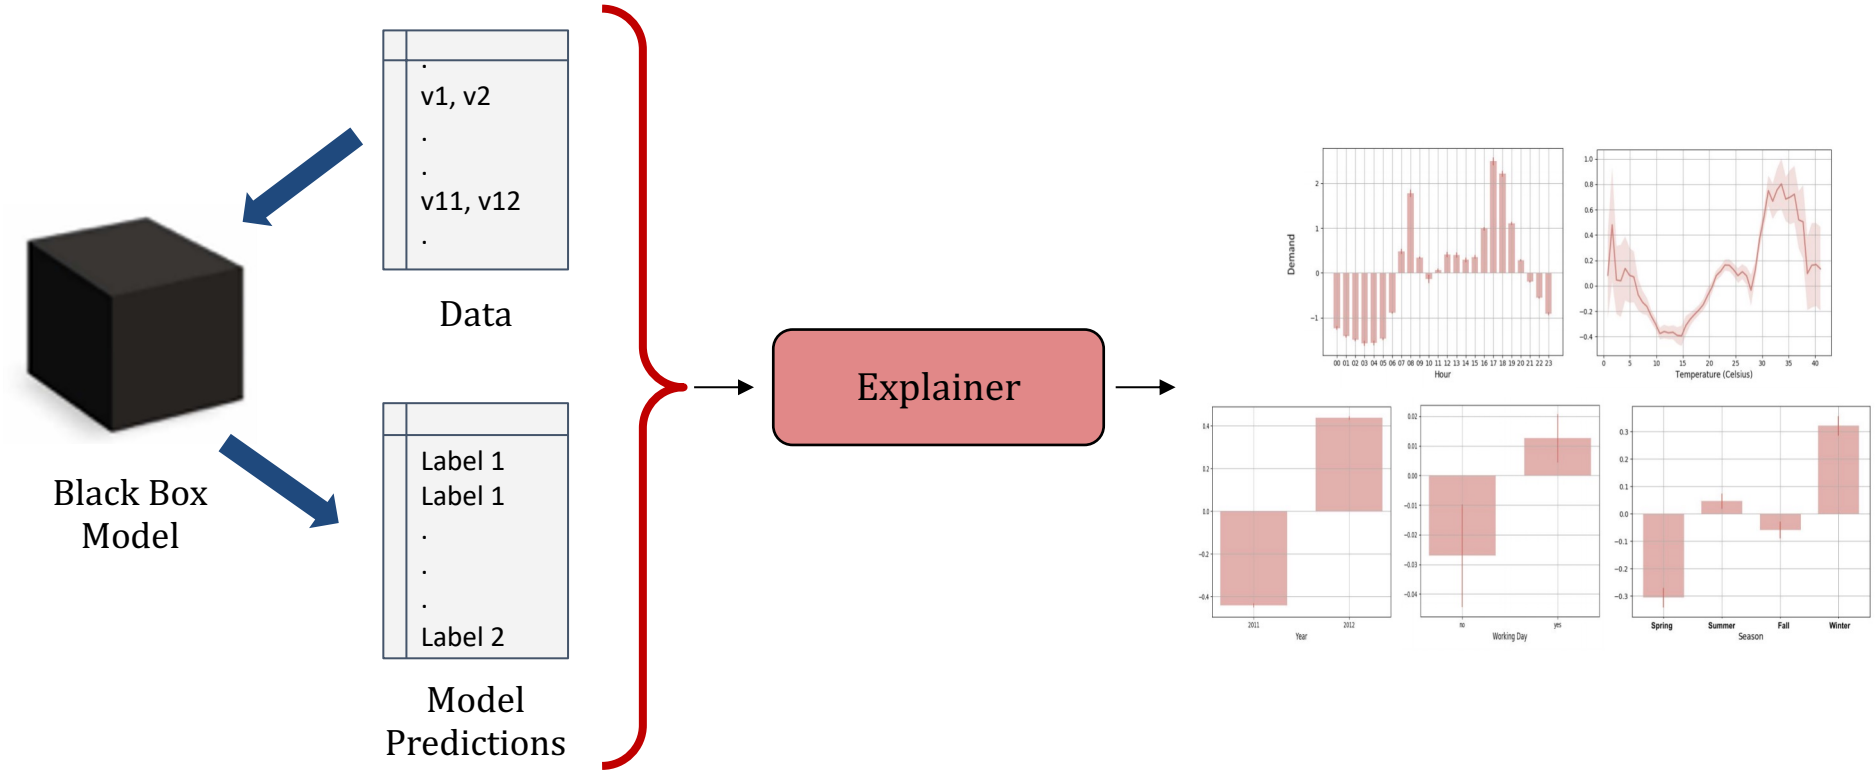

[Tan et. al., 2019]

Overview of explanation methods

- **Local explanation methods:**
  - Feature importance scoring
  - Integrated gradients
  - Prototype explanations
  - Counterfactuals

- **Global explanation methods:**
  - Collection of local explanations
  - Representation-based explanations
  - Model distillation

Outline for today's class

1. What is trustworthy AI/ML and why should I care?
2. Interpretability vs. explainability
3. Explaining AI/ML predictions
4. Case studies
   - Drug repurposing
   - Treatment recommendation

Rapid therapeutic innovation

- Pandemics demand safe and effective therapies developed at an unprecedented speed
- Traditional, iterative development, experimental and clinical testing, and approval of new drugs not feasible
- **Challenge:** How to compress years of work into months or even weeks through AI, automation, and new data resources?

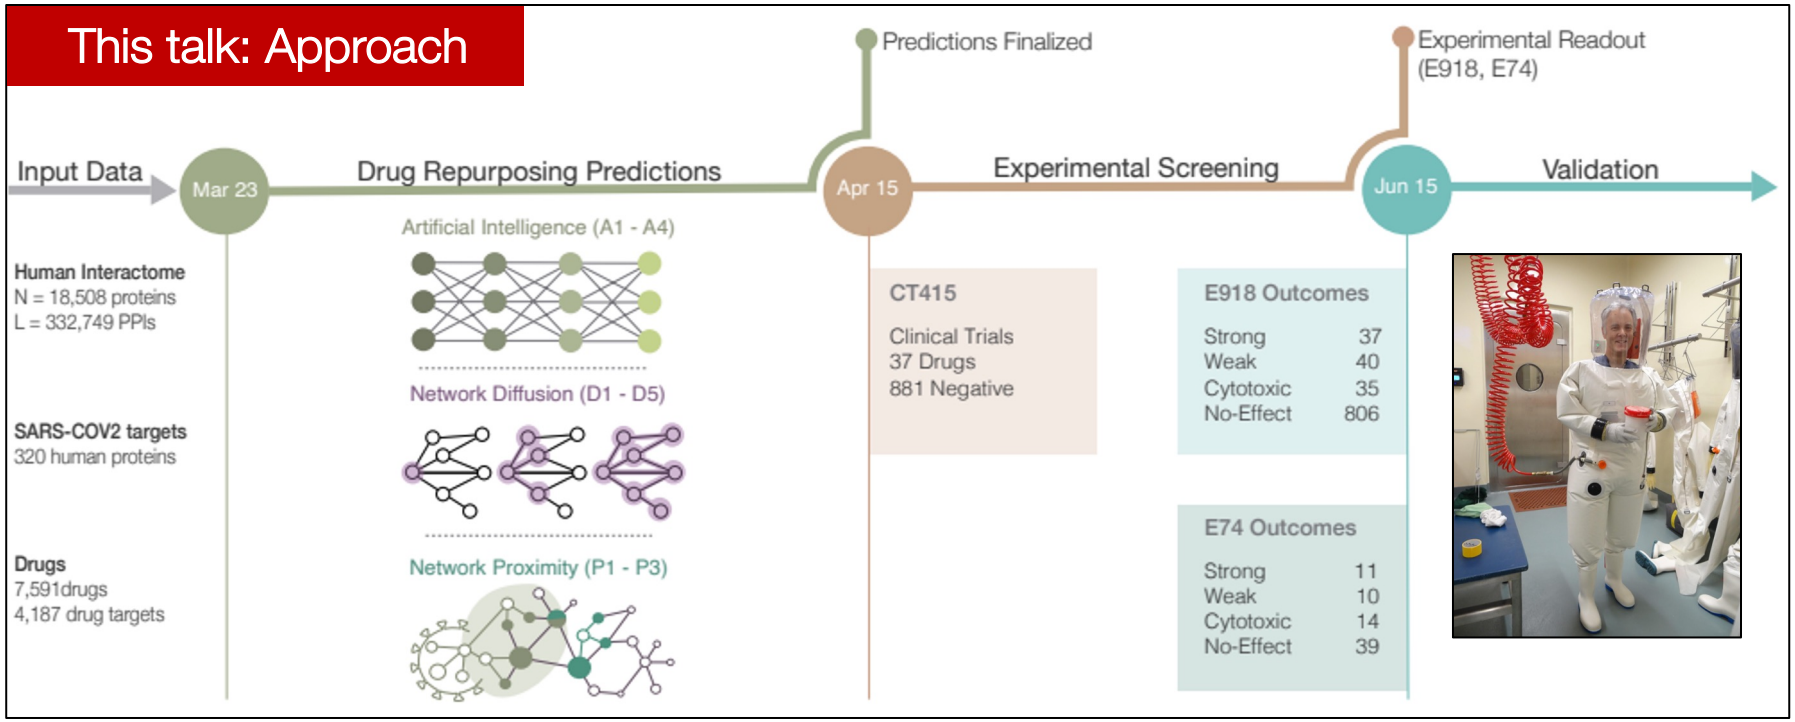

Network Medicine Framework for Identifying Drug Repurposing Opportunities for Covid-19, *PNAS*, 2021

Design therapies to target biological networks

Disease disrupts the normal behavior of genes. Drugs intervene against the disease by restoring the function of disrupted genes. **Goal:** What chemical compounds can intervene against disease?

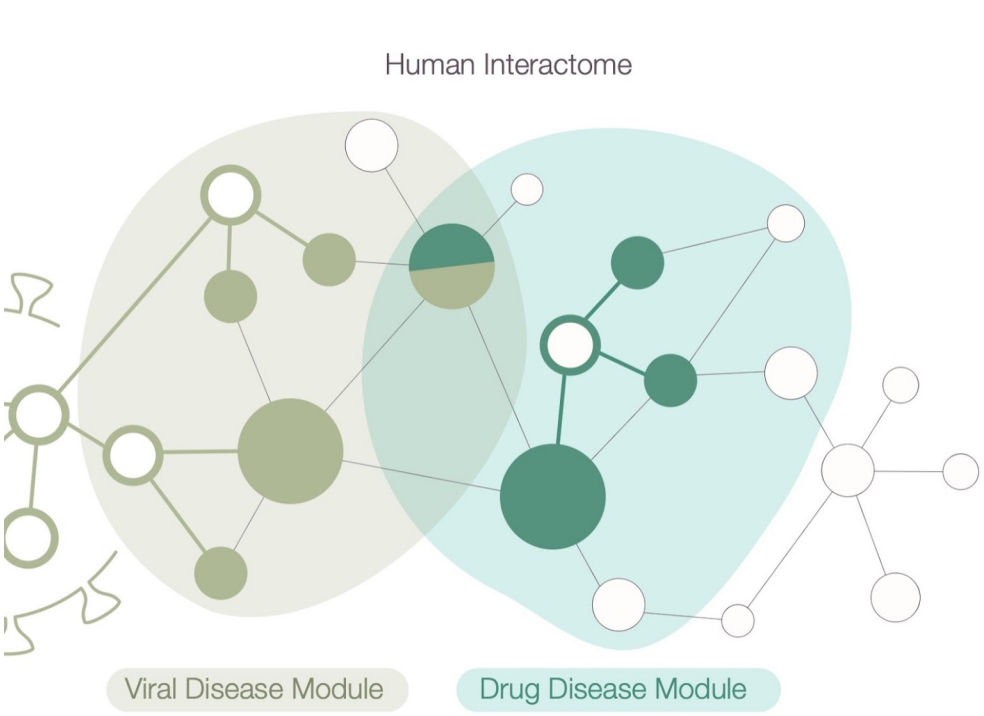

Human-Human
Protein-Protein Interaction

Viral-Human
Protein-Protein Interaction

Drug-Human
Protein-Protein Interaction

Network Medicine Framework for Identifying Drug Repurposing Opportunities for Covid-19, *PNAS*, 2021
Zero-shot prediction of therapeutic use with geometric deep learning and clinician centered design, medRxiv 2023

Dataset and experimental setup

- **COVID-19 repurposing knowledge graph:**
  - Human protein-protein interaction graph
  - Approved drugs and proteins that each drug targets
  - Diseases and proteins perturbed in each disease
  - Approved drug-disease treatments

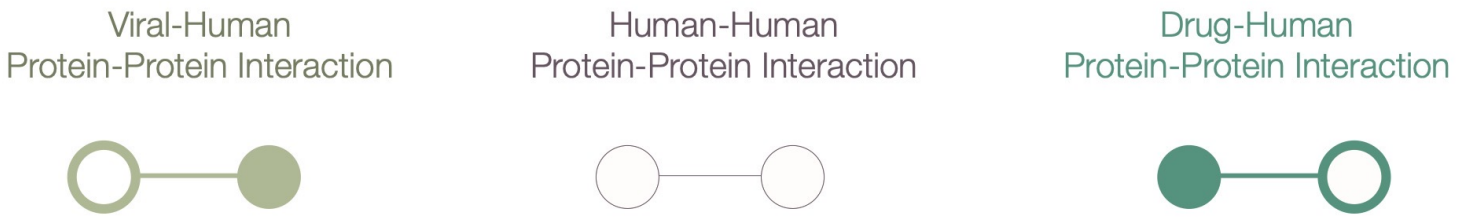

- **ML task:** Given approved drug-disease treatments, identify candidate treatments for COVID-19

Network Medicine Framework for Identifying Drug Repurposing Opportunities for Covid-19, *PNAS*, 2021

Approach: Graph ML model

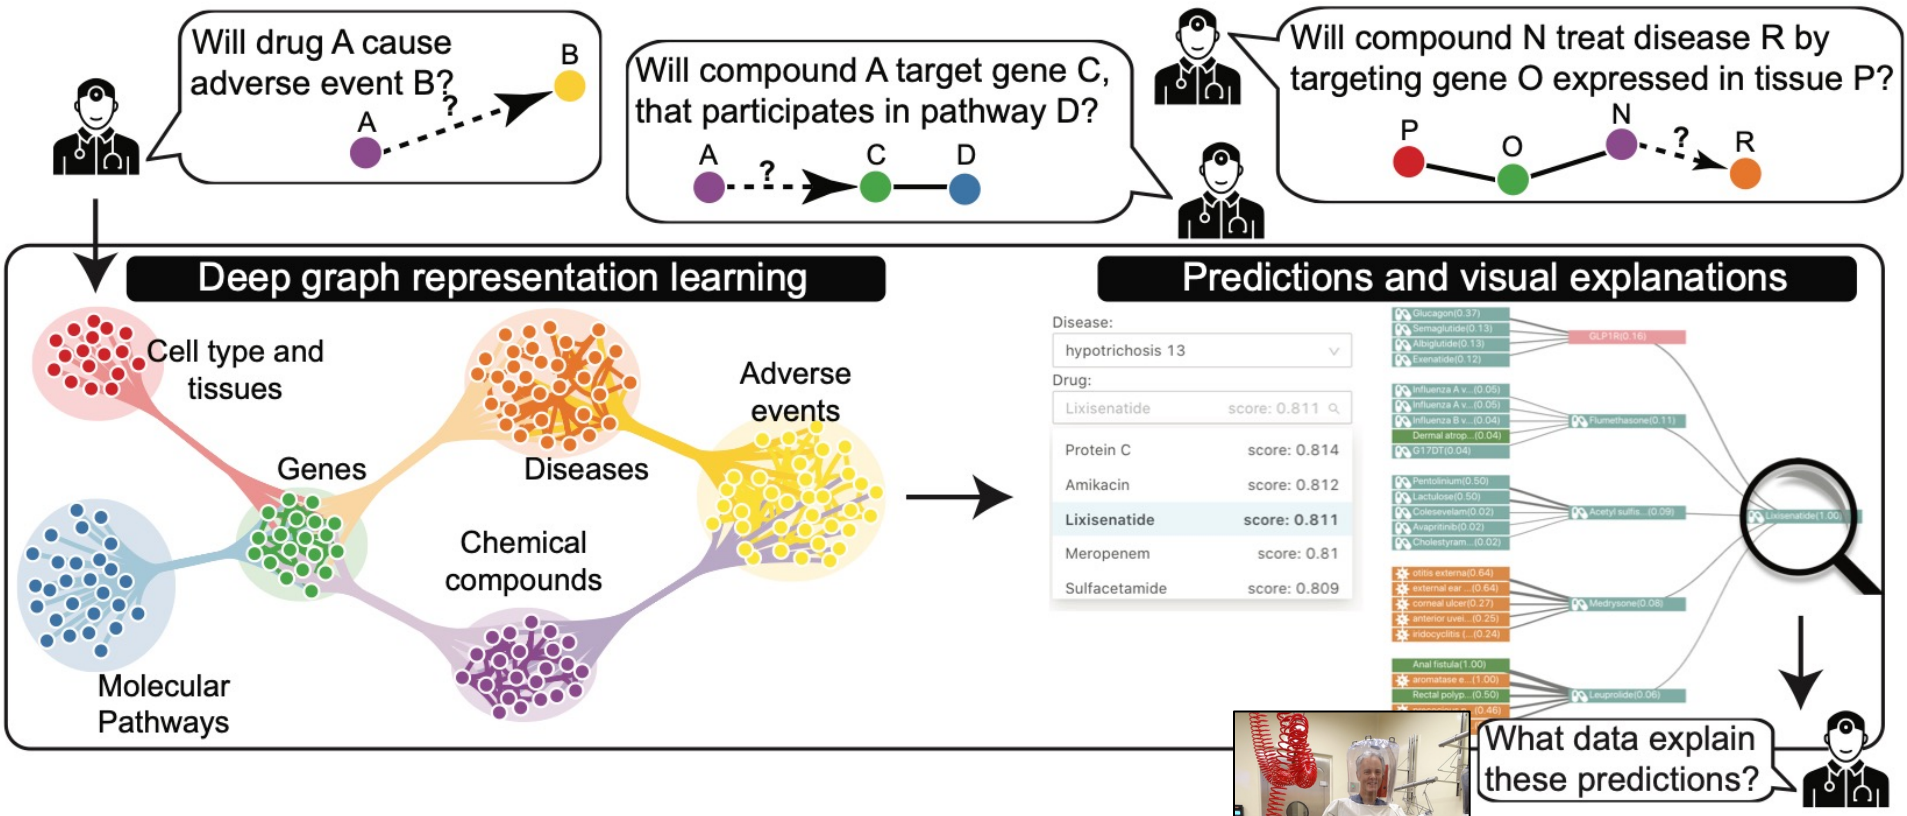

Approach: Graph ML model

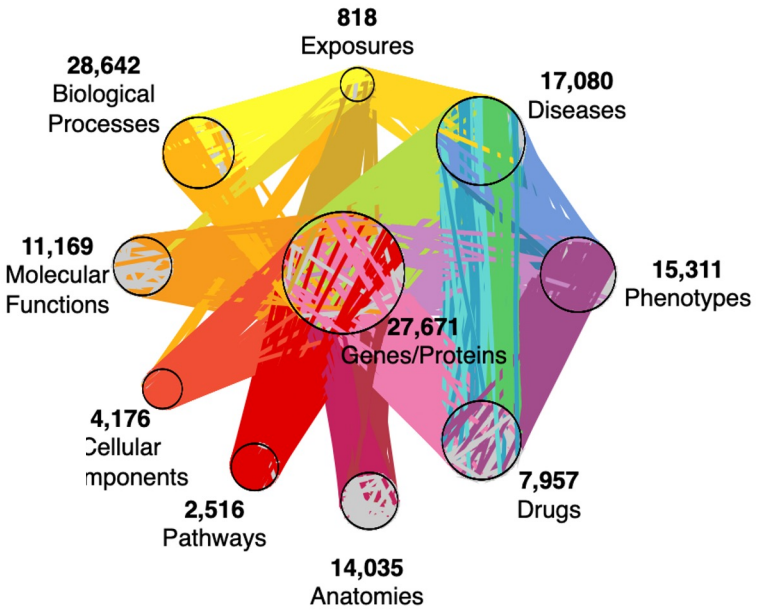

Model

Goal: What compounds/drugs can intervene against (treat) disease?

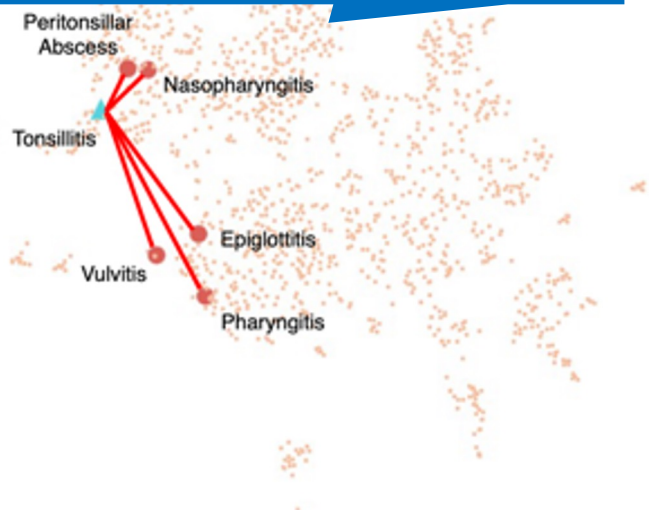

Results: COVID-19 repurposing

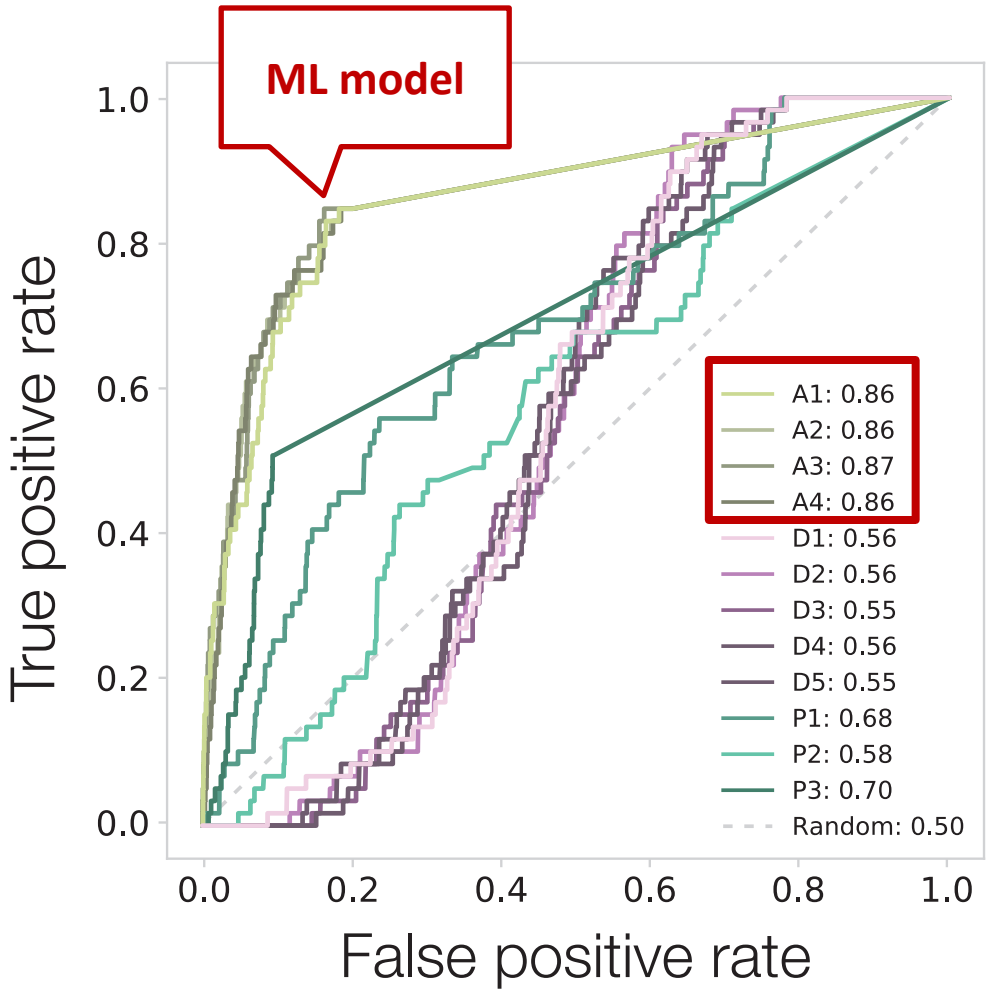

We test each method's ability to recover drugs currently in clinical trials for COVID-19 (67 drugs from ClinicalTrials.gov)

The best individual ROC curves are obtained by the GNN methods

The second-best performance is provided by the proximity P3. Close behind is P1 with AUC = 0.68 and AUC = 0.58

Diffusion methods offer ROC between 0.55-0.56

Network Medicine Framework for Identifying Drug Repurposing Opportunities for Covid-19, *PNAS*, 2021

Results: Experimental screening

National Emerging Infectious Diseases Laboratories (NEIDL)

CRank	Drug Name
1	Ritonavir
2	Isoniazid
3	Troleandomycin
4	Cilostazol
5	Chloroquine
6	Rifabutin
7	Flutamide
8	Dexamethasone
9	Rifaximin
10	Azelastine
11	Crizotinib

17	Celecoxib
18	Betamethasone
19	Prednisolone
20	Mifepristone
21	Budesonide
22	Prednisone
23	Oxiconazole
24	Megestrol acetate
25	Idelalisib
26	Econazole

Predicted lists of drugs

New algorithms:
Prioritizing Network Communities, *Nature Communications* 2018
Subgraph Neural Networks, *NeurIPS* 2020
Graph Meta Learning via Local Subgraphs, *NeurIPS* 2020

Results: 918 compounds screened for their efficacy against SARS-CoV-2 in VeroE6 & human cells:

- We screened in human cells the top-ranked drugs, obtaining a 62% success rate, in contrast to the 0.8% hit rate of nonguided screenings
- This is an **order of magnitude higher hit rate** among top 100 drugs than **alternative approach**

Network Medicine Framework for Identifying Drug Repurposing Opportunities for Covid-19, *PNAS*, 2021

Results: Predicting therapeutic use

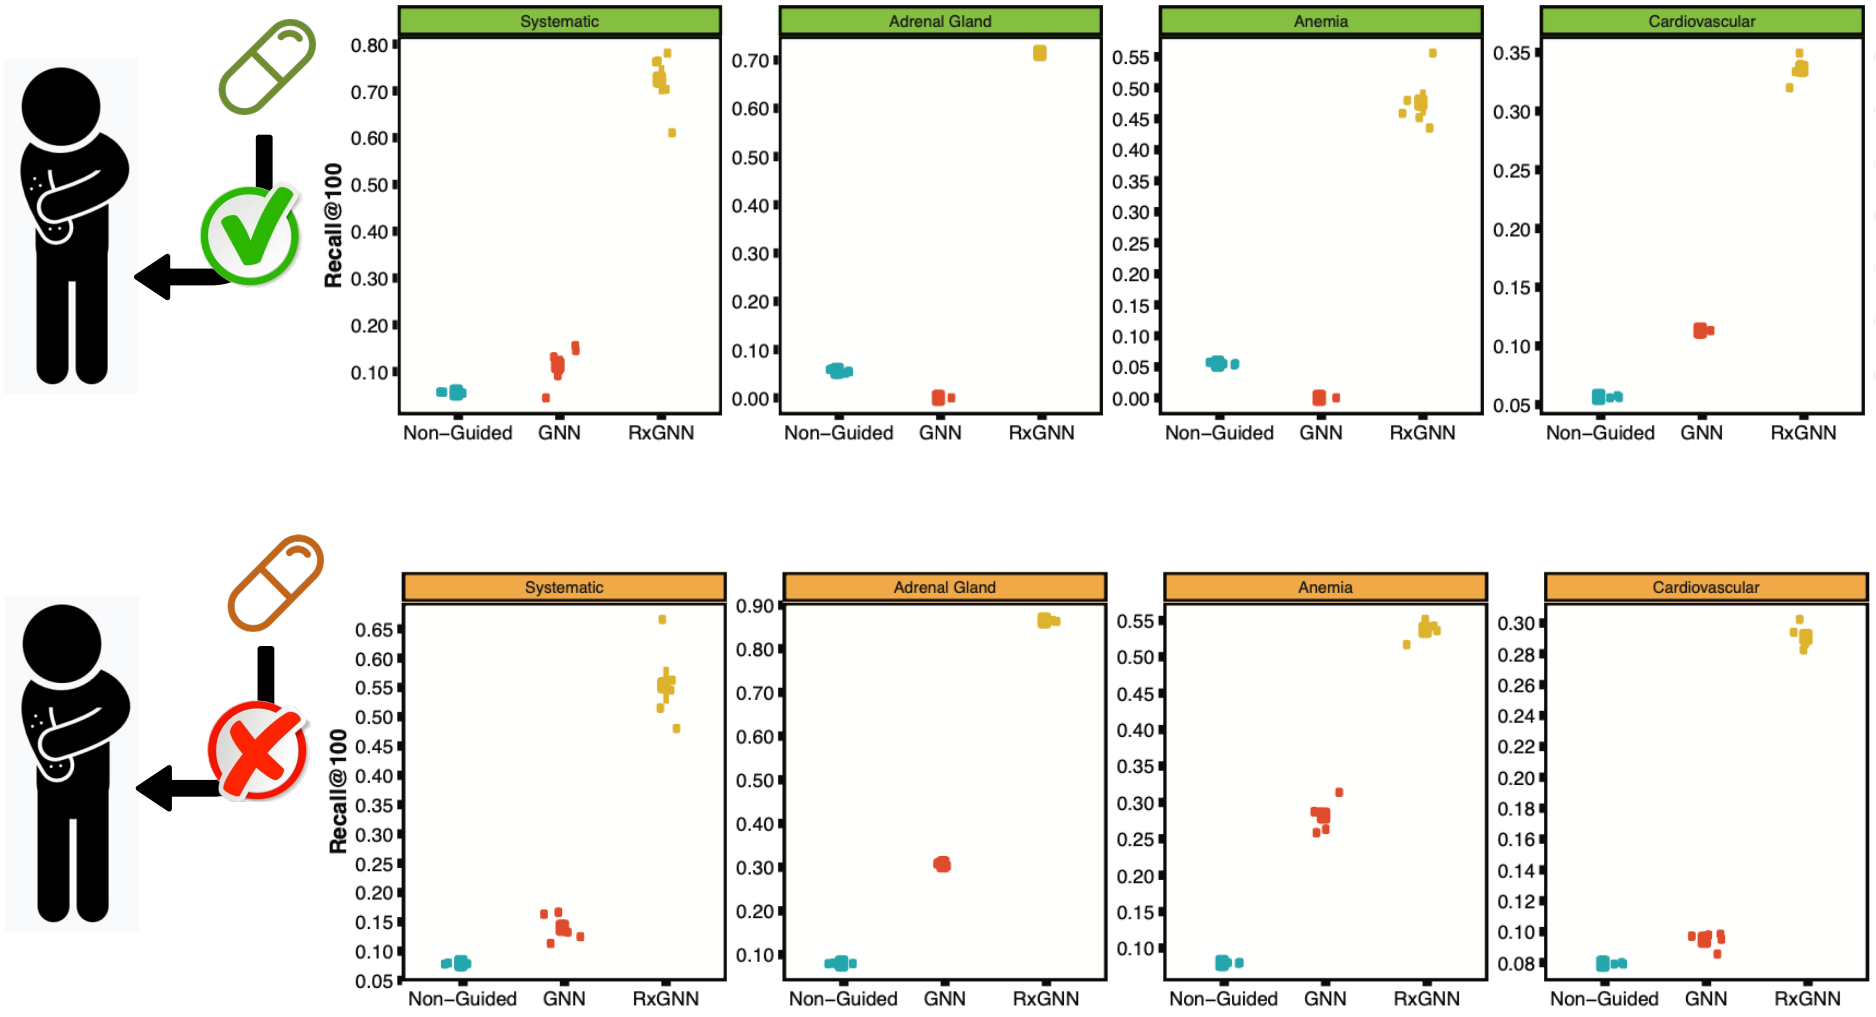

Explaining machine predictions

Key idea:

- Summarize where in the data the model "looks" for evidence for its prediction
- Find a small subgraph most influential for the prediction

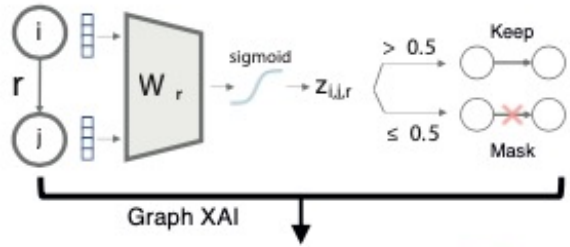

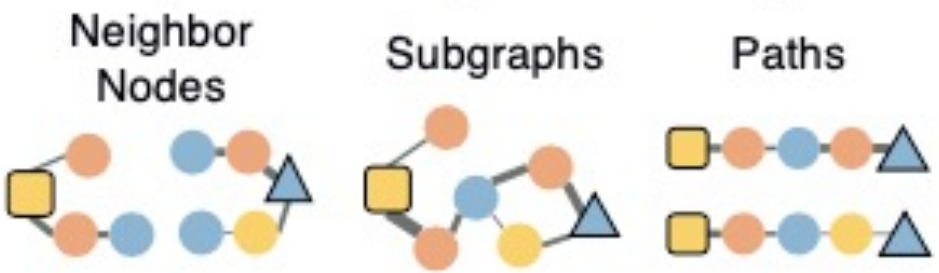

Interpreting Graph Neural Networks for NLP With Differentiable Edge Masking, *ICLR* 2021
GNNExplainer: Generating Explanations for Graph Neural Networks, *NeurIPS* 2019

GNNExplainer: Key idea

- Input: Given prediction $f(x)$ for node/link x
- Output: Explanation, a small subgraph M_x together with a small subset of node features:
  - M_x is most influential for prediction $f(x)$
- Approach: Optimize mask M_x in a post-hoc manner
  - Intuition: If removing v from the graph strongly decreases the probability of prediction $\Rightarrow$ v is a good counterfactual explanation for the prediction

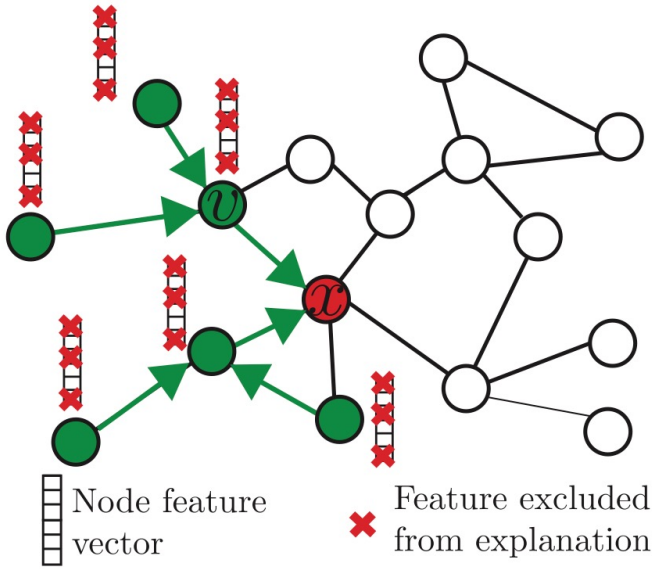

GNN Explainer: Generating Explanations for Graph Neural Networks, *NeurIPS* 2019

Explanations: Network drugs

"What is the disease treatment mechanism for drugs with **positive experimental outcomes**?"

76/77 predicted drugs with positive experimental outcomes do not directly bind to SARS-CoV-2 targets:

- Instead, the drugs rely on **network-based actions** and cannot be identified by alternative, non-network methods

Direct target drugs (D1 and D2)

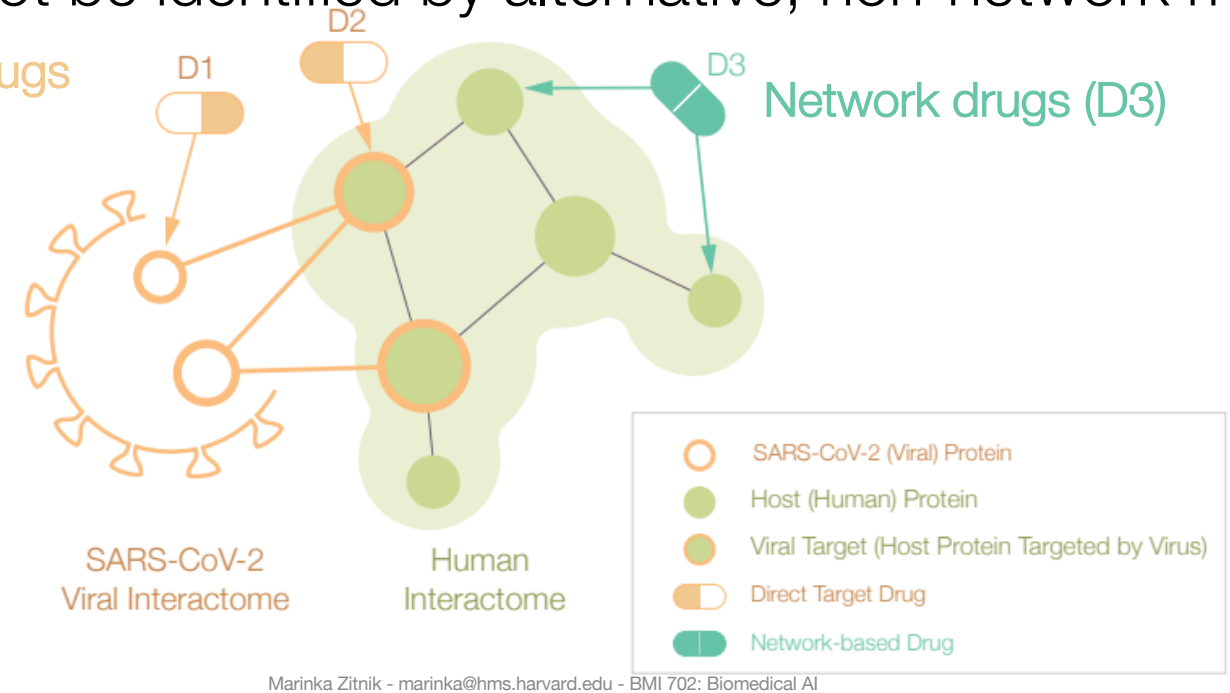

AI-clinician collaboration

"Will clozapine treat unipolar depression? What is the disease treatment mechanism?"

Mount Sinai

Support scientists in interacting with AI predictions and interpreting conclusions of AI analyses

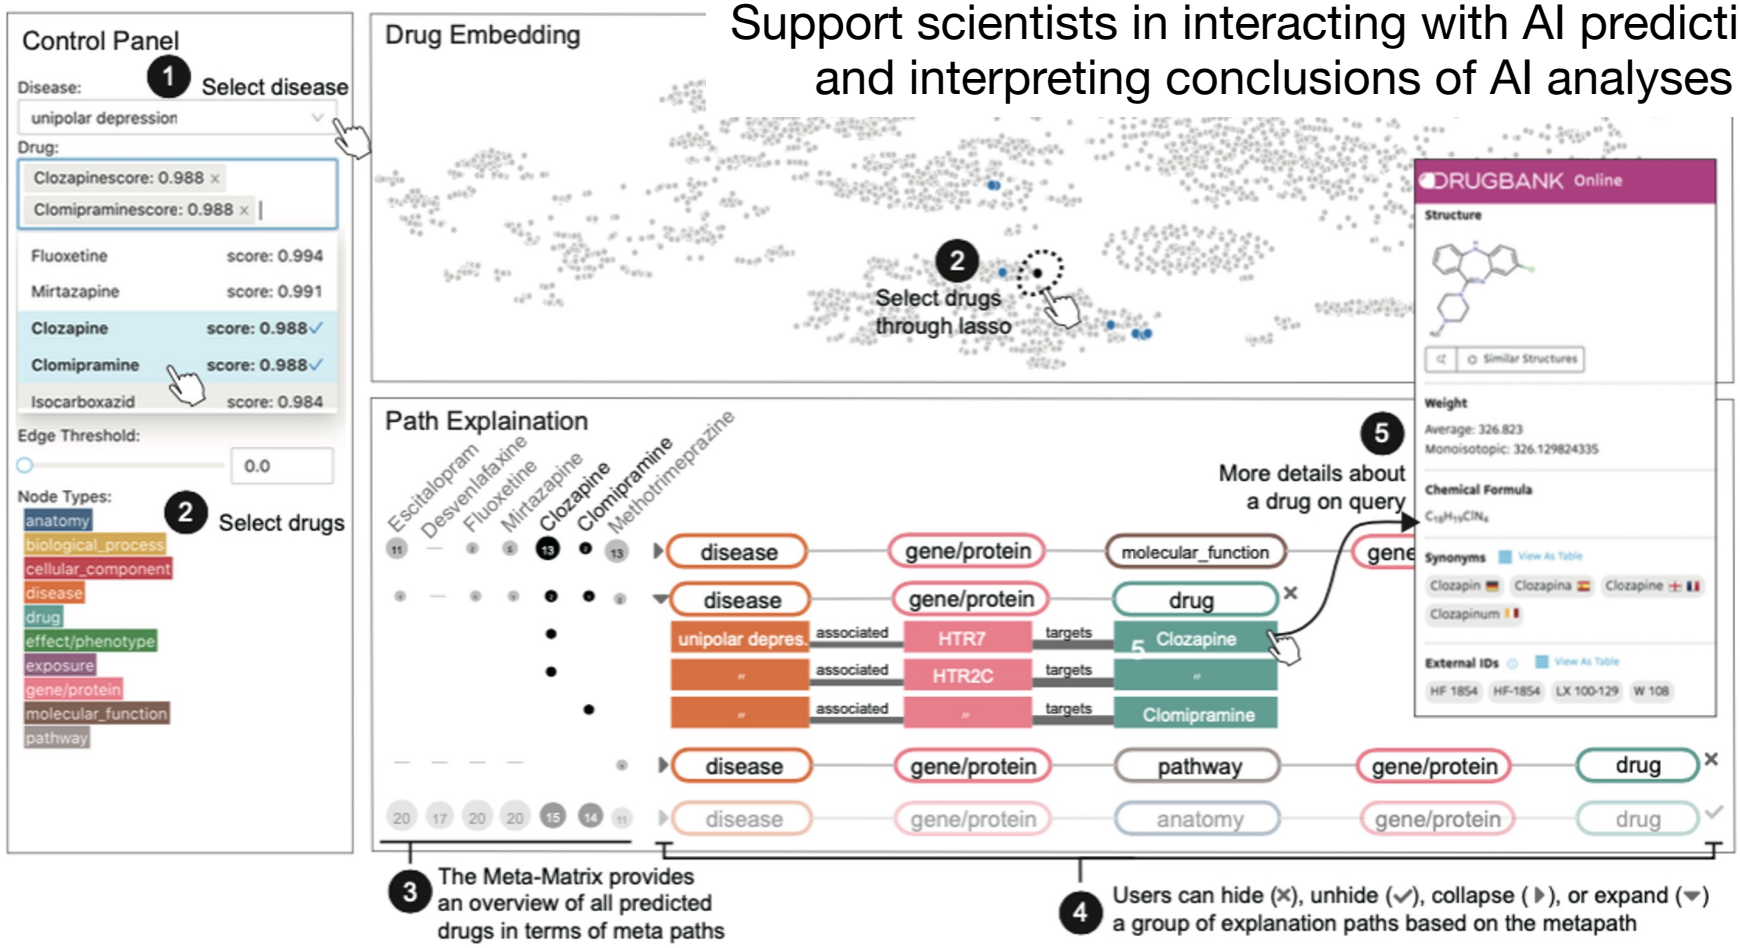

Probing GNN Explainers: A Rigorous Theoretical and Empirical Analysis of GNN Explanation Methods, *AISTATS* 2022

Extending the Nested Model for User-Centric XAI: A Design Study on GNN-based Drug Repurposing, *IEEE VIS* 2022 (**Best Paper Award**)

Identification of Disease Treatment Mechanisms through the Multiscale Interactome, Nature Communications 2021

Clinician-centric study

Compared to a no-explanation baseline in terms of **user answer accuracy**, **exploration time**, **user confidence**, and **user agreement** across a spectrum of usability questions

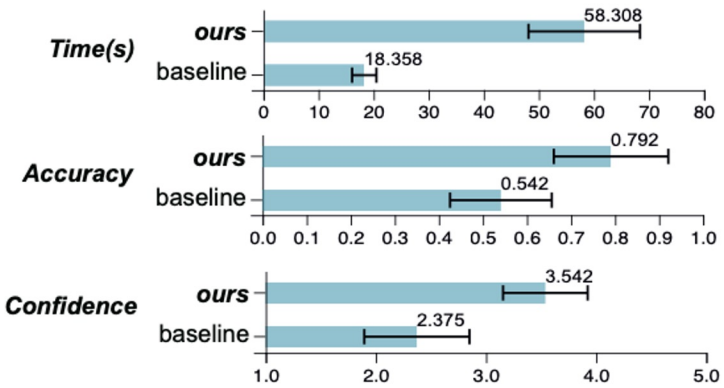

Error bars indicate the 95% confidence intervals

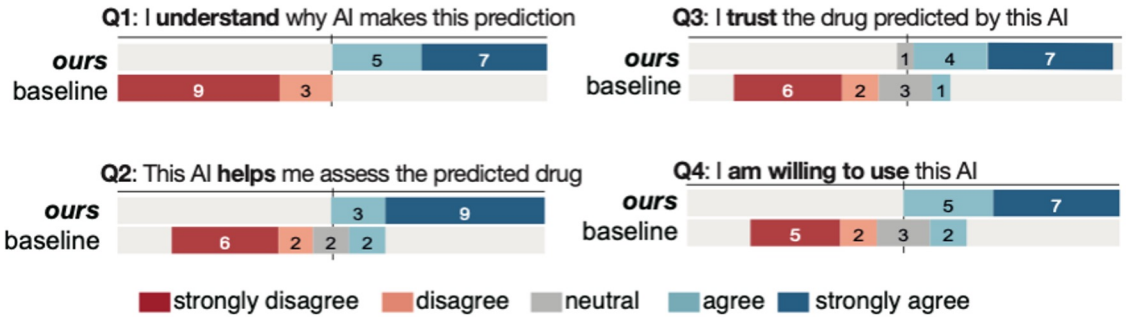

Agree scores are placed to the right, disagree to the left

Practical and ethical challenges

Q: Are decision-makers benefitting from explanations?
A: (Mixed) evidence of real-world benefit

Q: How are explanations calibrating trust in AI?
A: Explanations can be used to manipulate & miscalibrate trust

Q: How are explanations calibrating perceptions of fairness?
A: Explanations can be used to change fairness perceptions

Q: Can adversaries fool explanation algorithms & hence users?
A: Adversaries can easily obfuscate true model behavior

Outline for today's class

1. What is trustworthy AI/ML and why should I care?
2. Interpretability vs. explainability
3. Explaining AI/ML predictions
4. Case studies
   - Drug repurposing
   - Treatment recommendation